

Hadoop. Apache Spark.


СУБД или Hadoop?

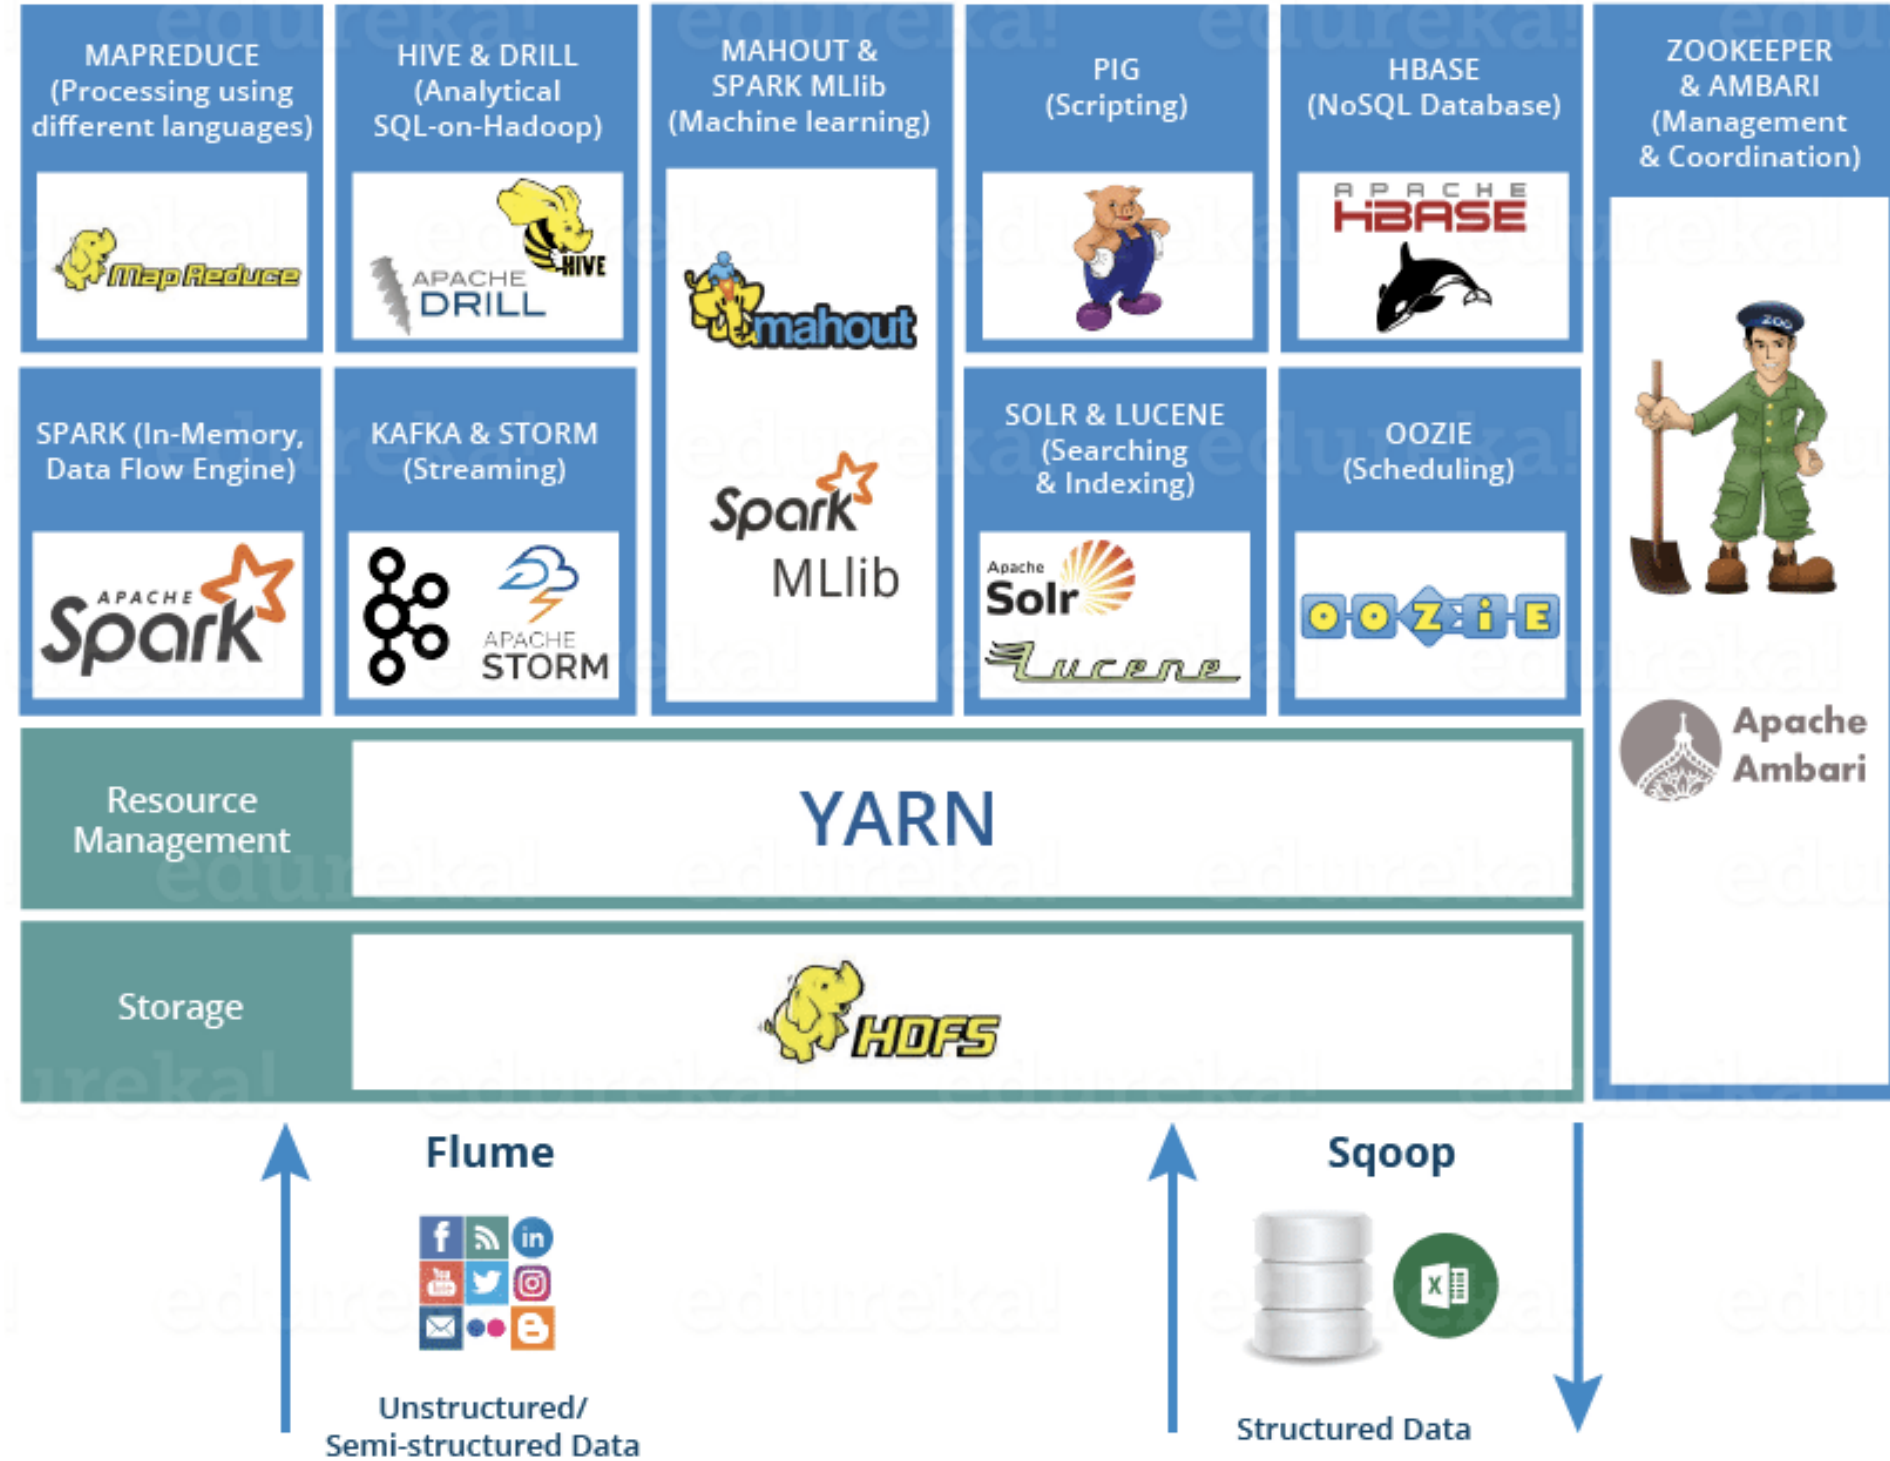
История Hadoop

Что такое Hadoop и Apache Spark?

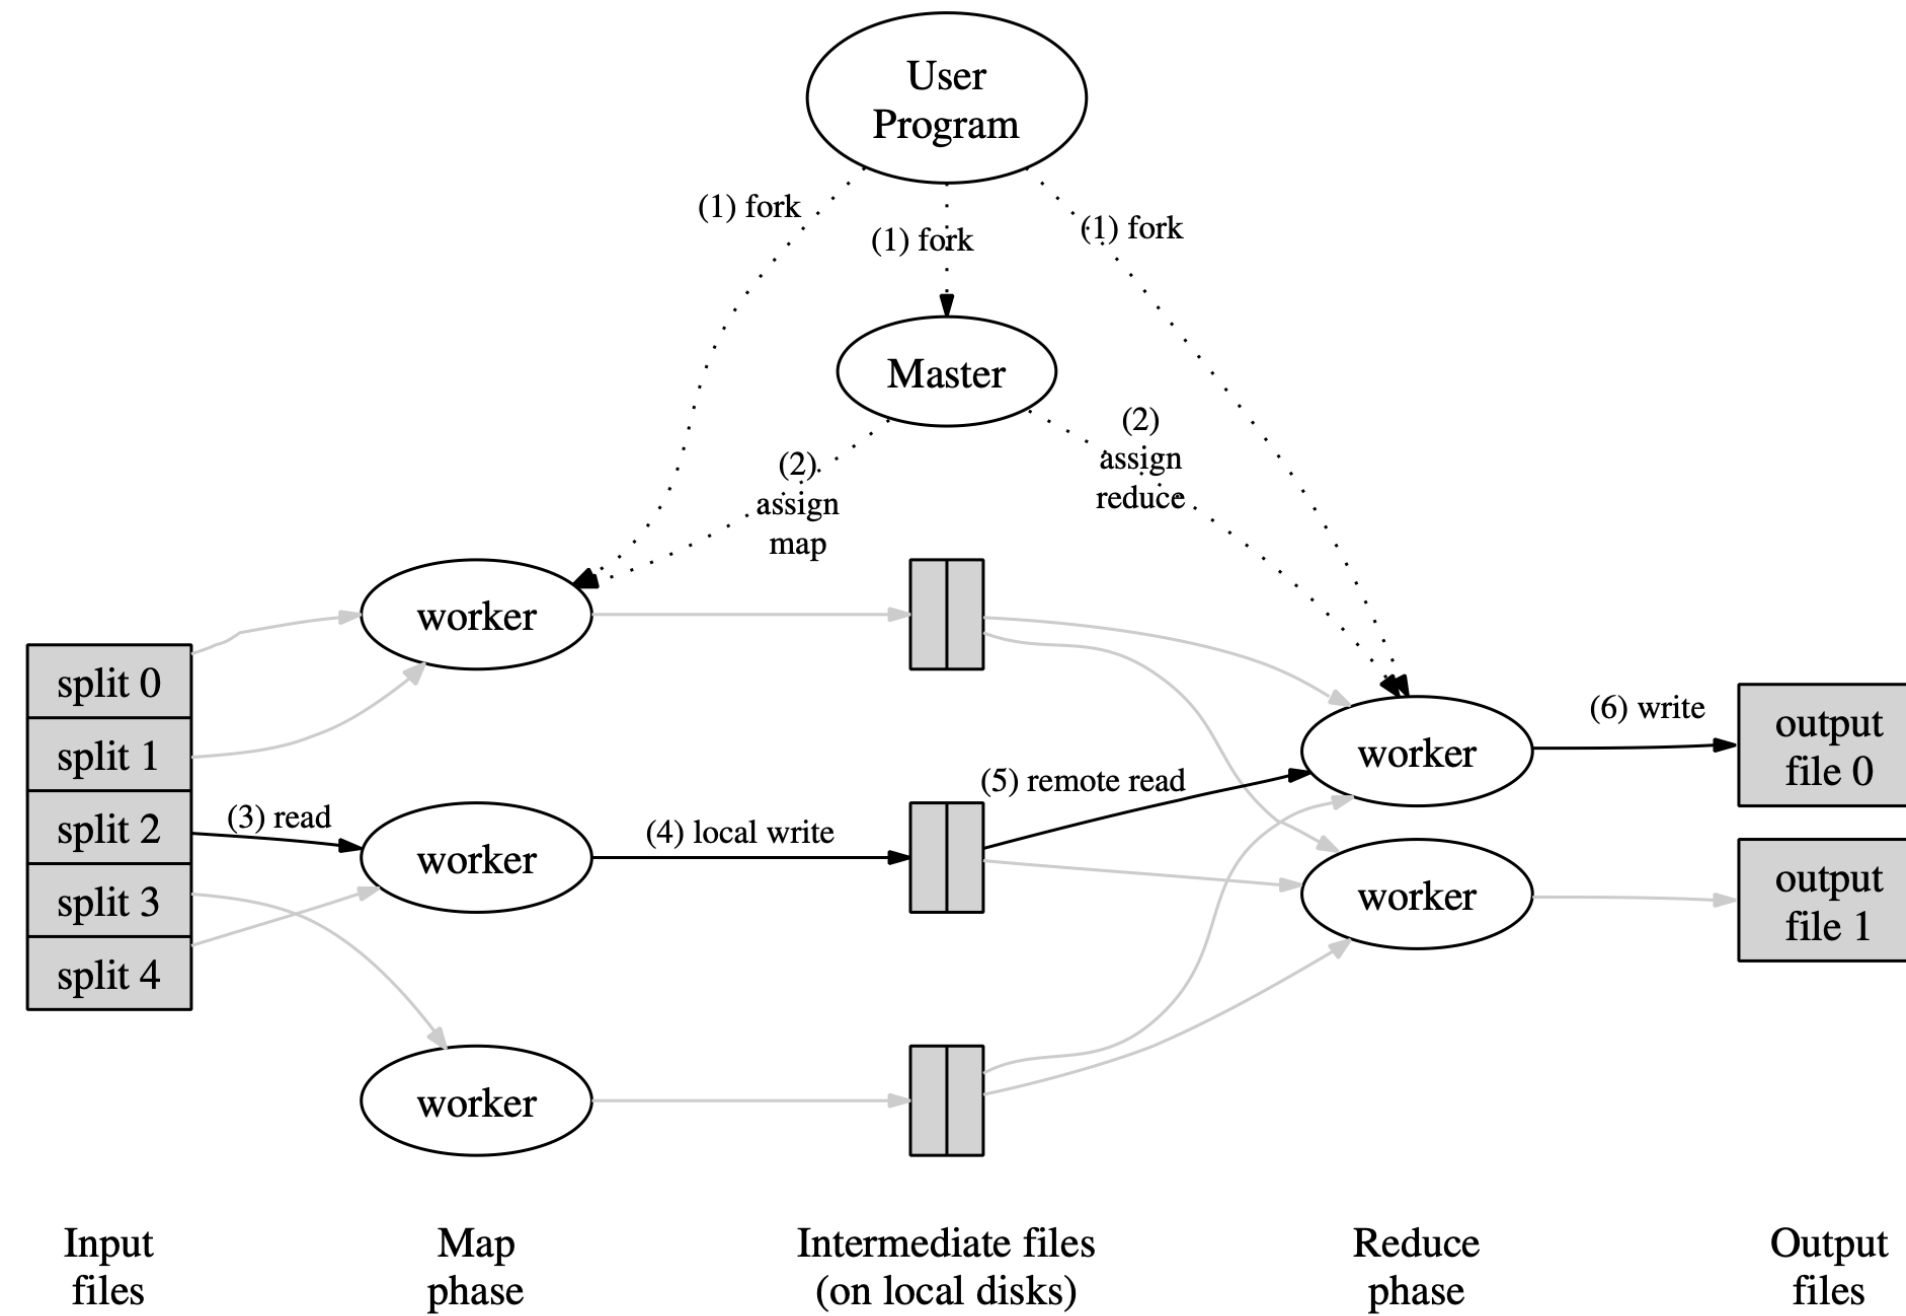
Hadoop можно рассматривать как экосистему, включающую в себя набор программ для параллельной обработки больших данных. Разработан в рамках вычислительной парадигмы **MapReduce**, согласно которой приложение разделяется на большое количество одинаковых элементарных заданий, выполнимых на узлах кластера и естественным образом сводимых в конечный результат.

Hadoop Ecosystem

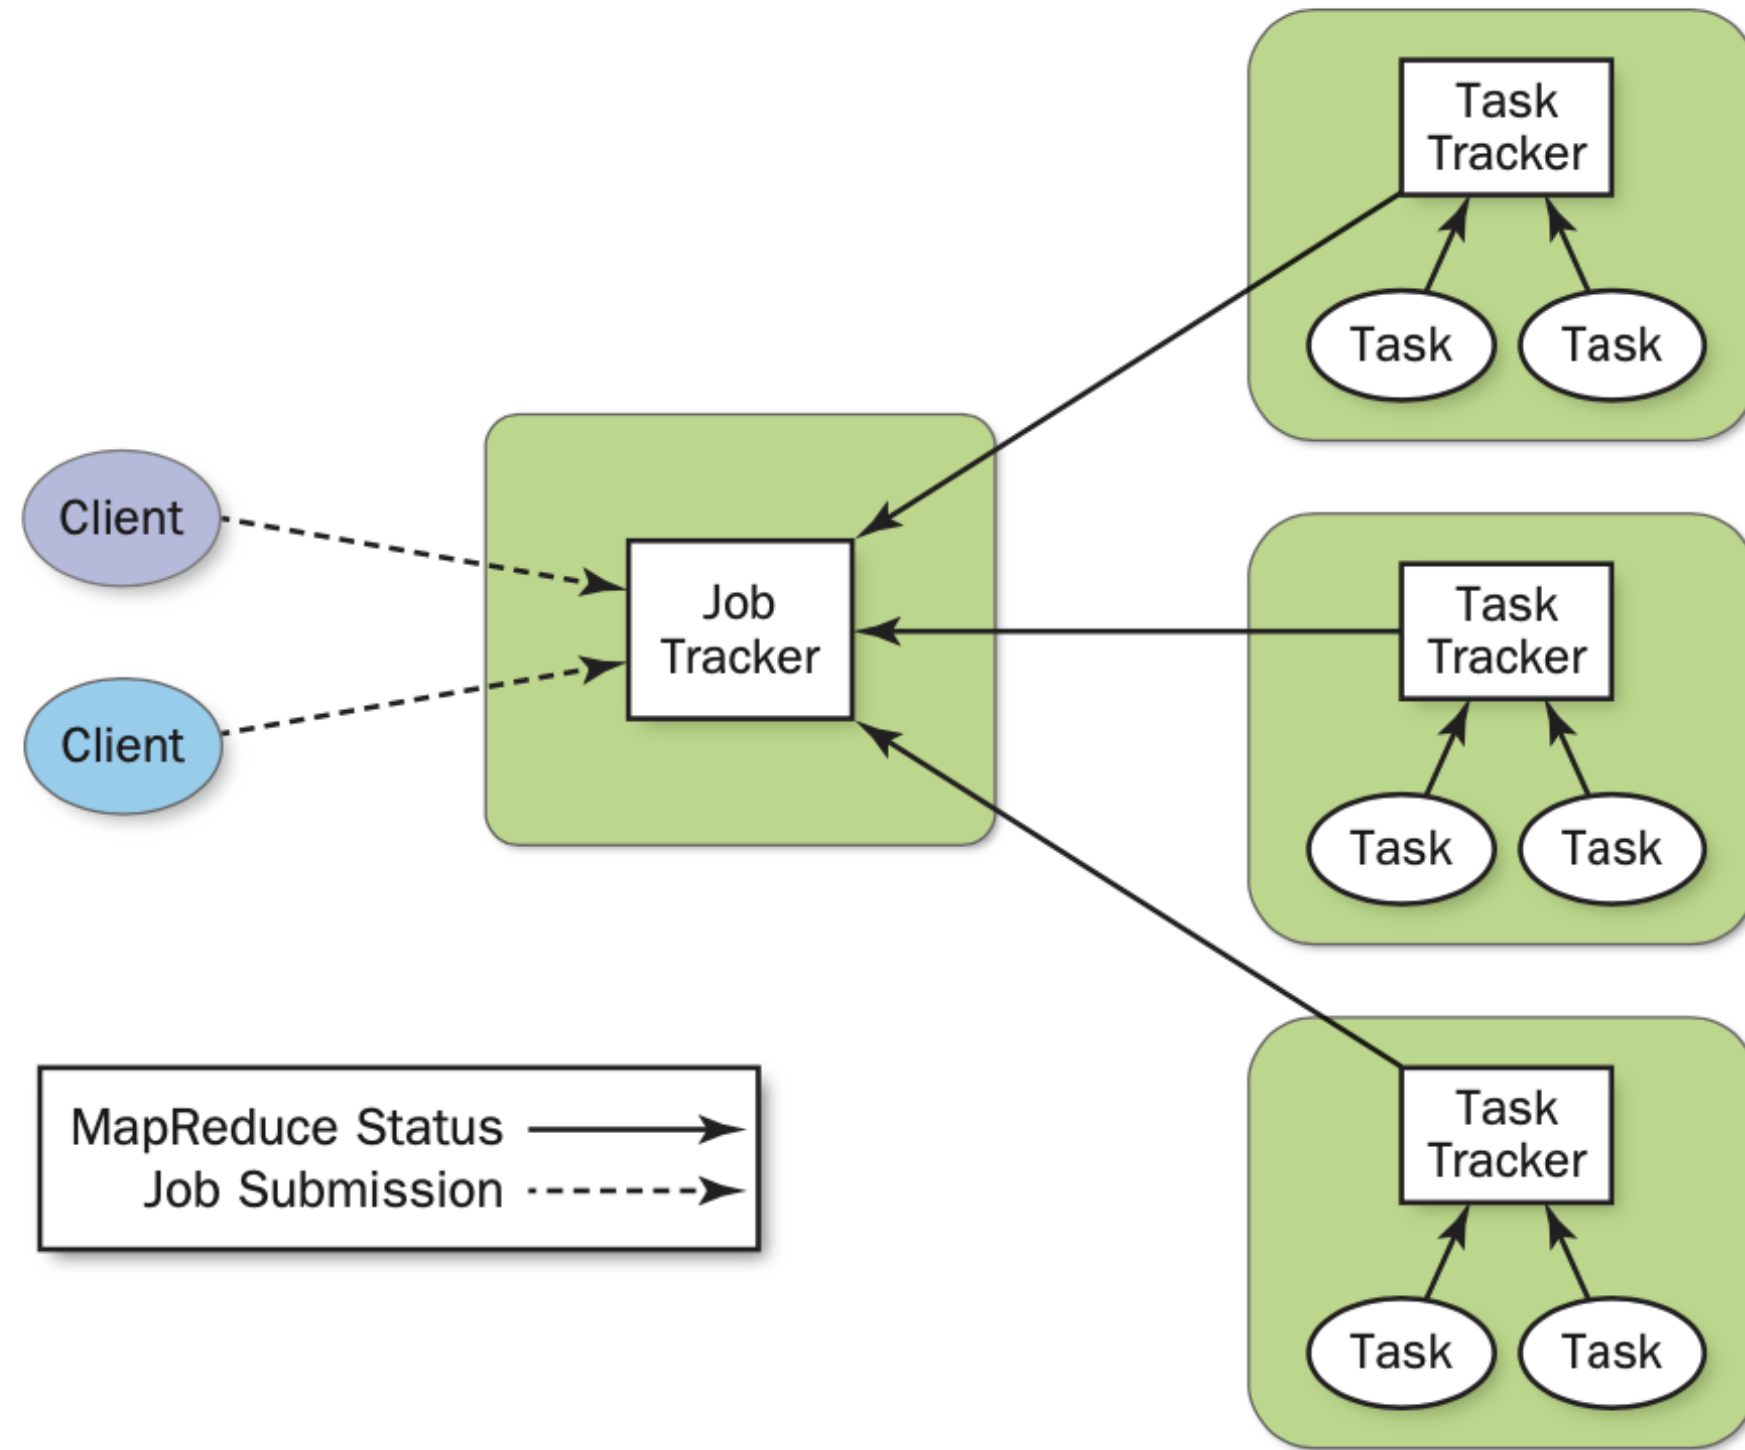






MapReduce



Jeffrey Dean, Sanjay Ghemawat. MapReduce: simplified data processing on large clusters. Proceedings of OSDI, 2004.



Функция Map





```
function map(String name, String document):  
  // name: document name  
  // document: document contents  
  for each word w in document:  
    emit (w, 1)
```

Mapper.py

```
#!/usr/bin/env python
import sys

for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for word in words:
        print '%s\t%s' % (word, 1)
```

Функция Reduce



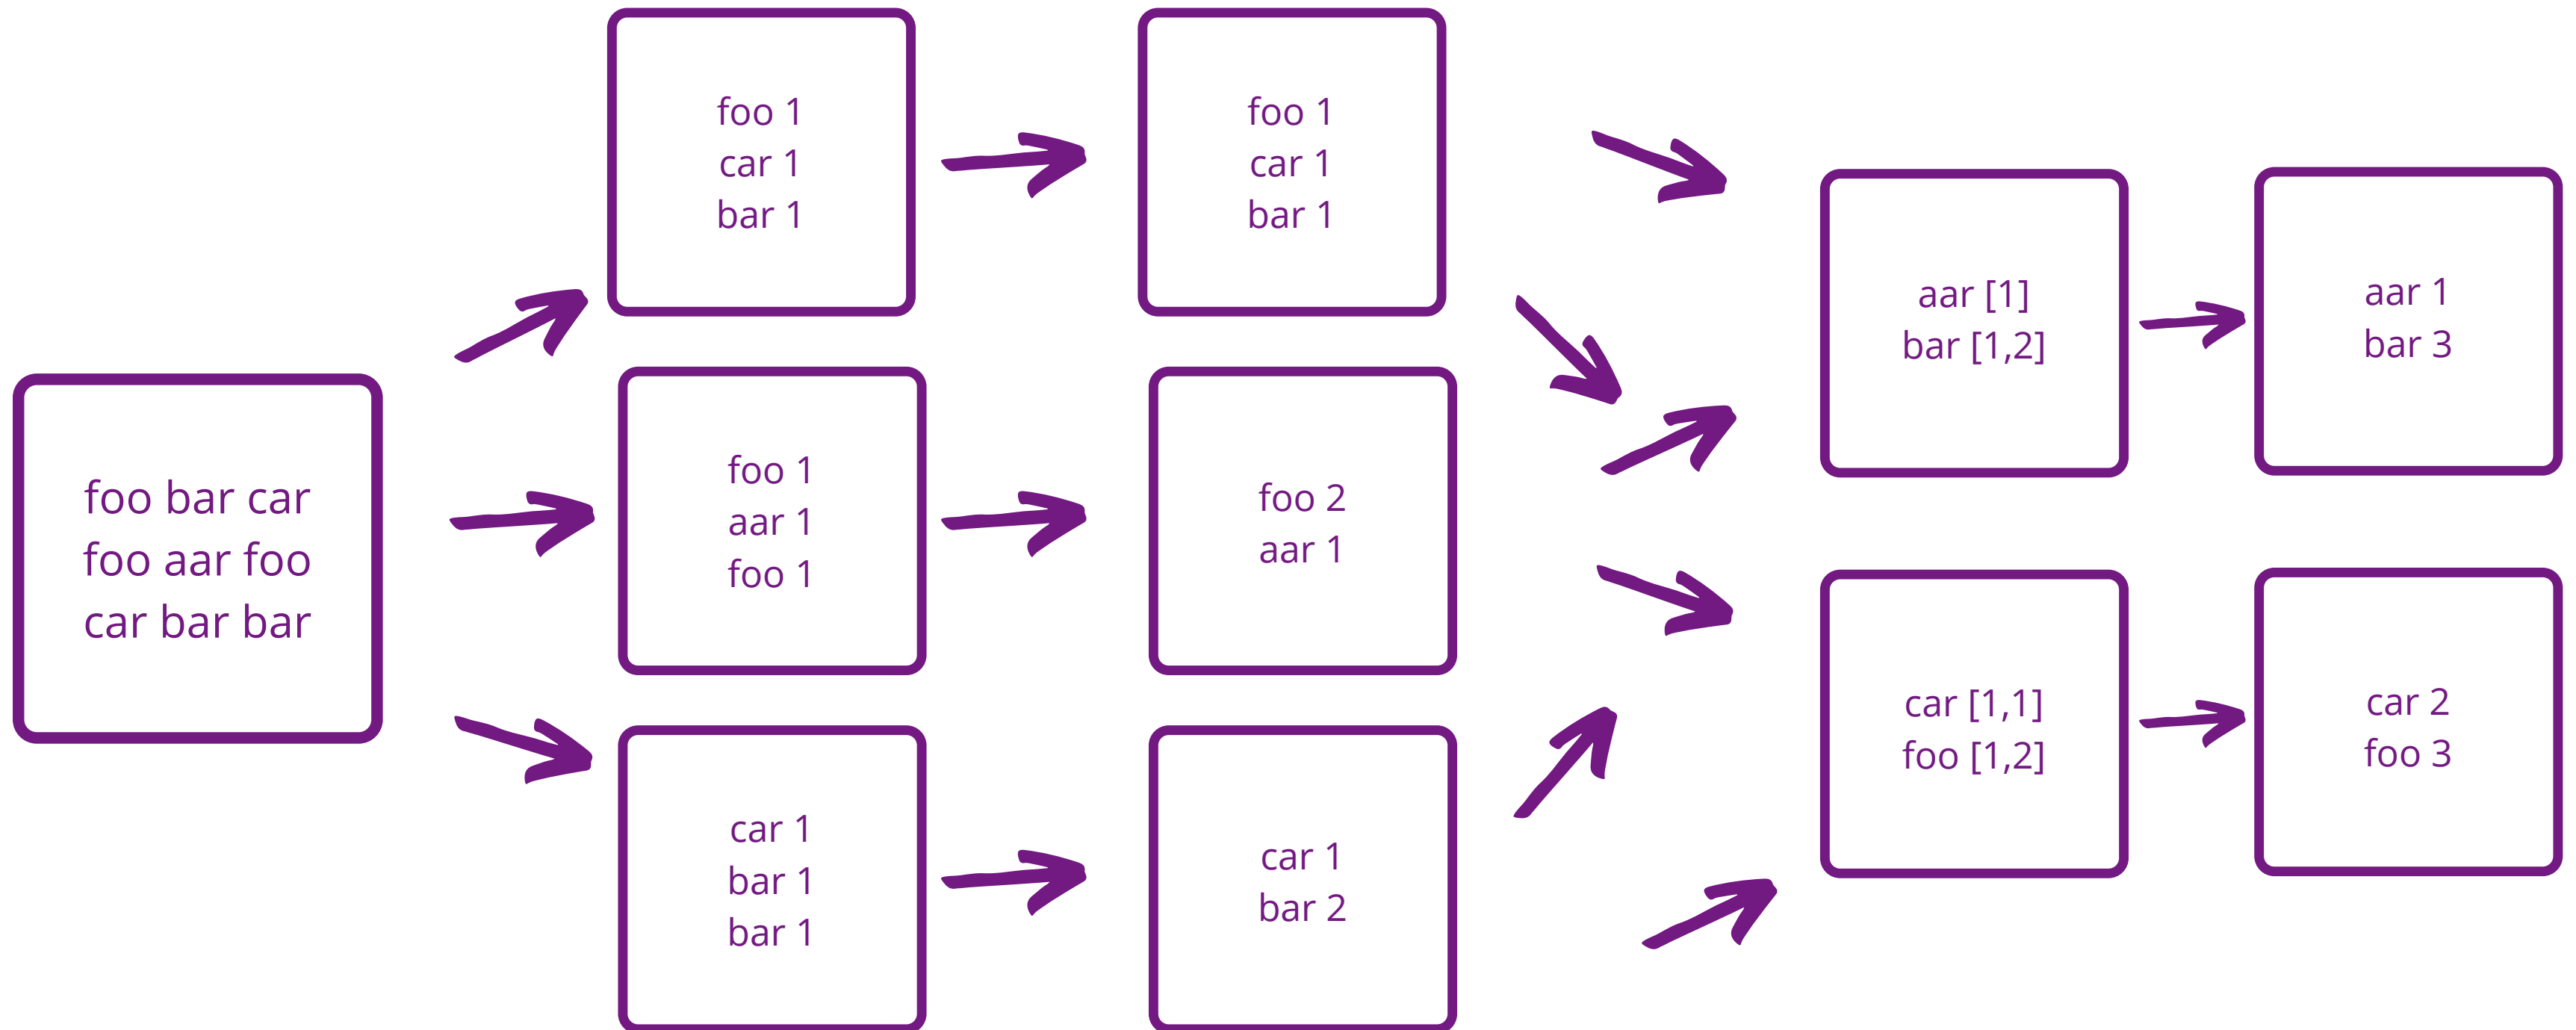
```
function reduce(String word, Iterator partialCounts):  
    // word: a word  
    // partialCounts: a list of aggregated partial counts  
    sum = 0  
    for each pc in partialCounts:  
        sum += pc  
        emit (word, sum)
```

Reducer.py

```
#!/usr/bin/env python
from operator import itemgetter
import sys

current_word = None
current_count = 0
word = None
for line in sys.stdin:
    line = line.strip()
    word, count = line.split('\t', 1)
    try:
        count = int(count)
    except ValueError:
        continue
    if current_word == word:
        current_count += count
    else:
        if current_word:
            print '%s\t%s' % (current_word, current_count)
            current_count = count
            current_word = word

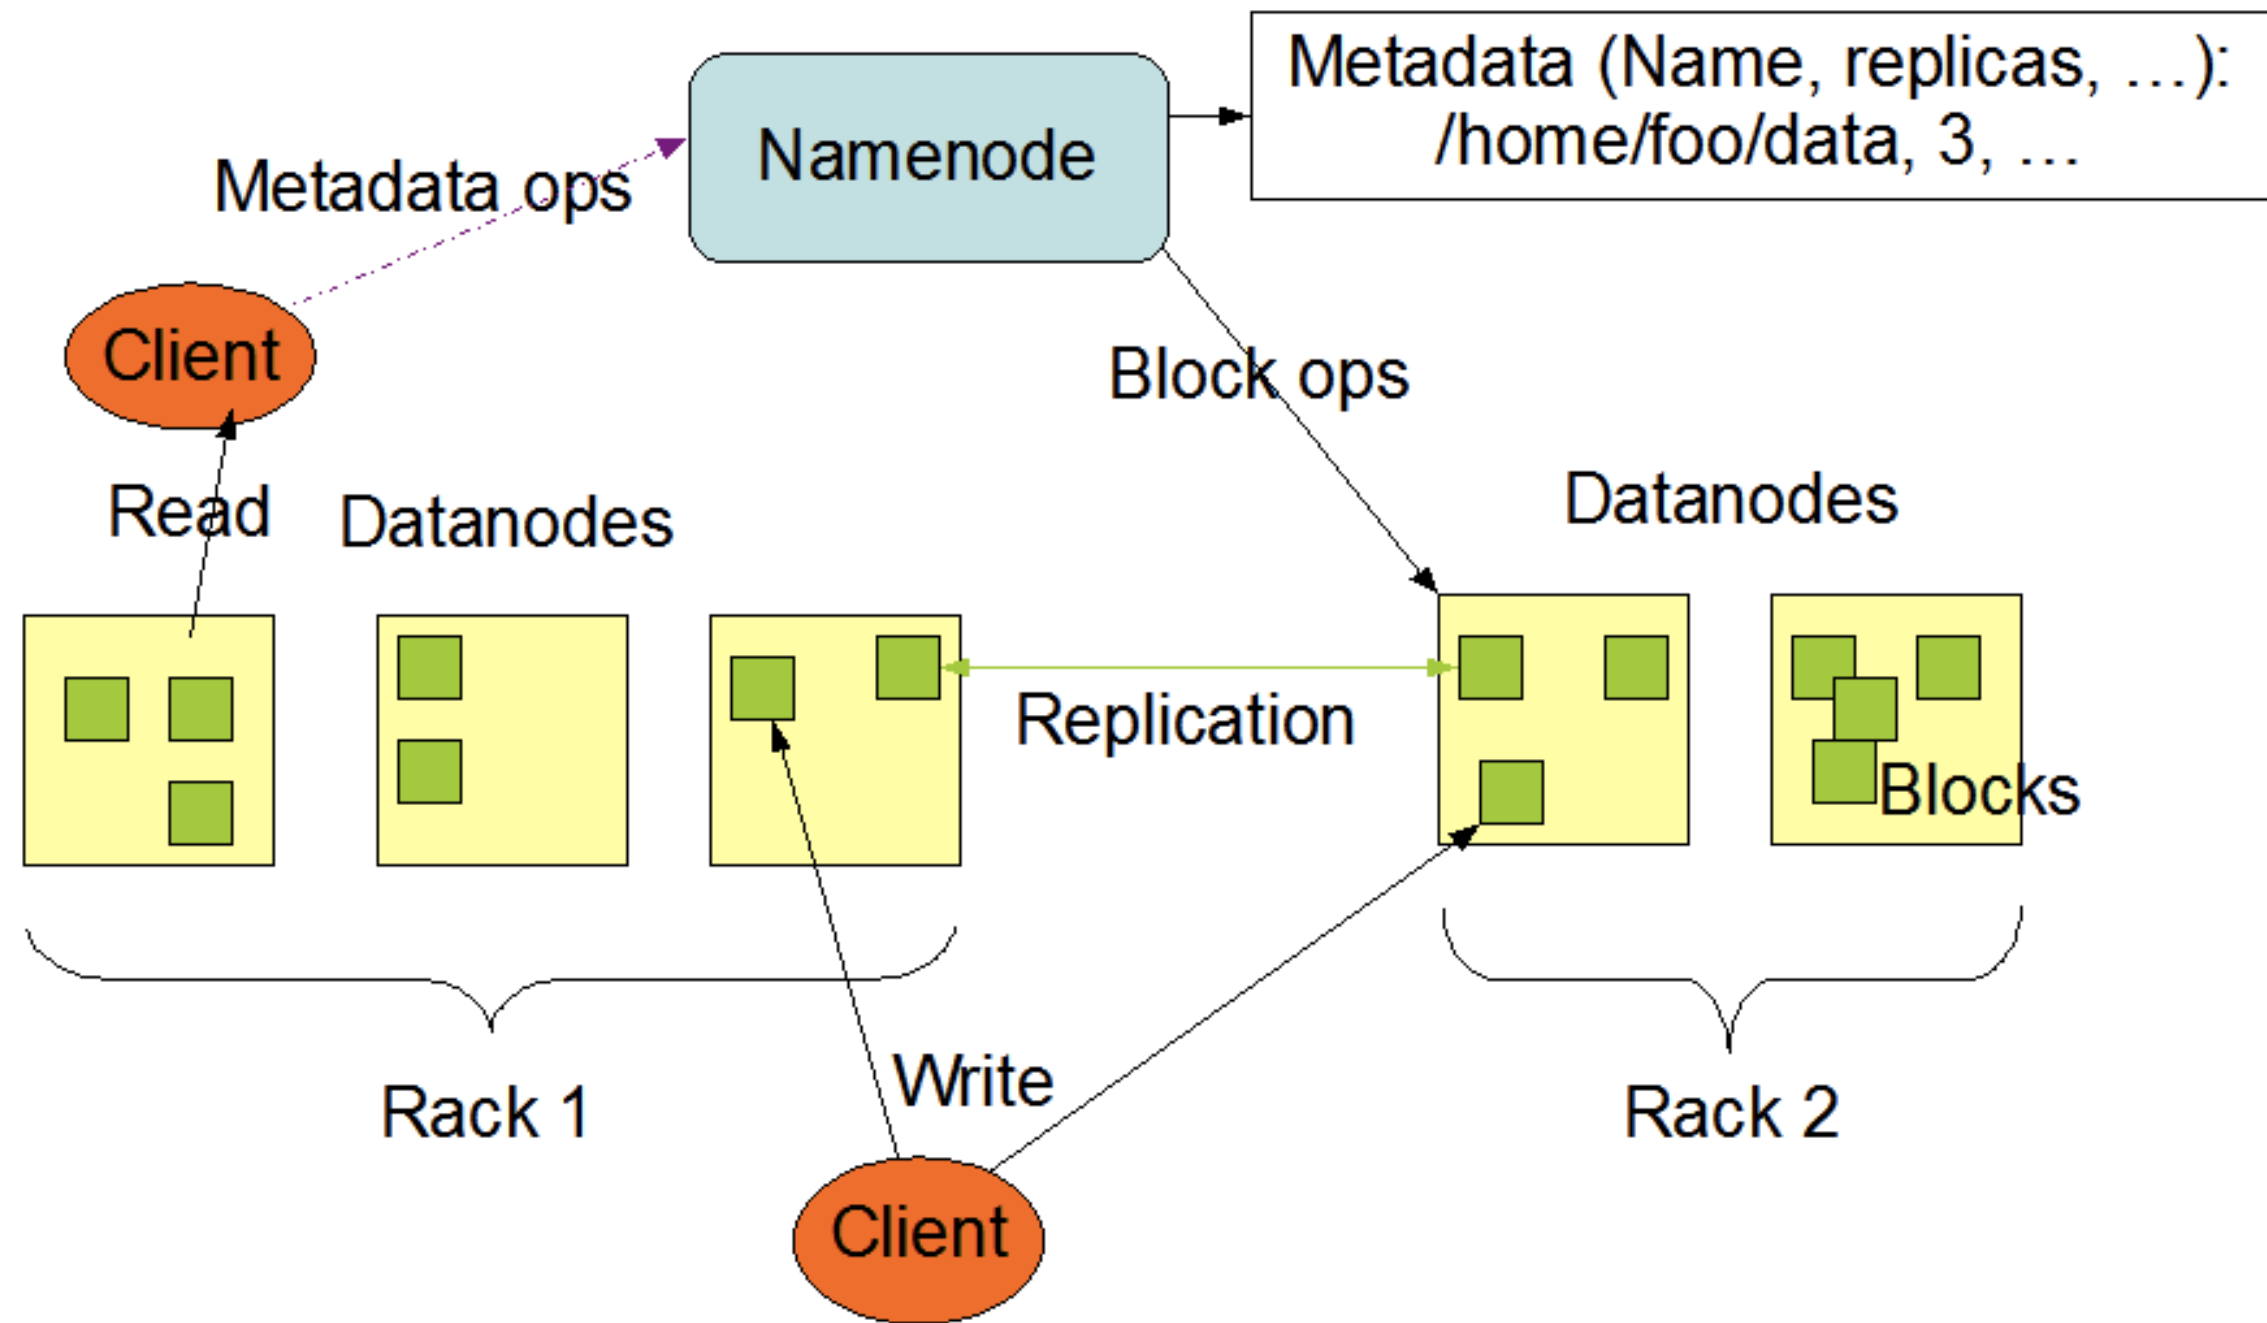
if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```

HDFS

- Hardware Failure
- Streaming Data Access
- Large Data Sets
- Simple Coherency Model
- “Moving Computation is Cheaper than Moving Data”
- Portability Across Heterogeneous Hardware and Software Platforms

HDFS Architecture

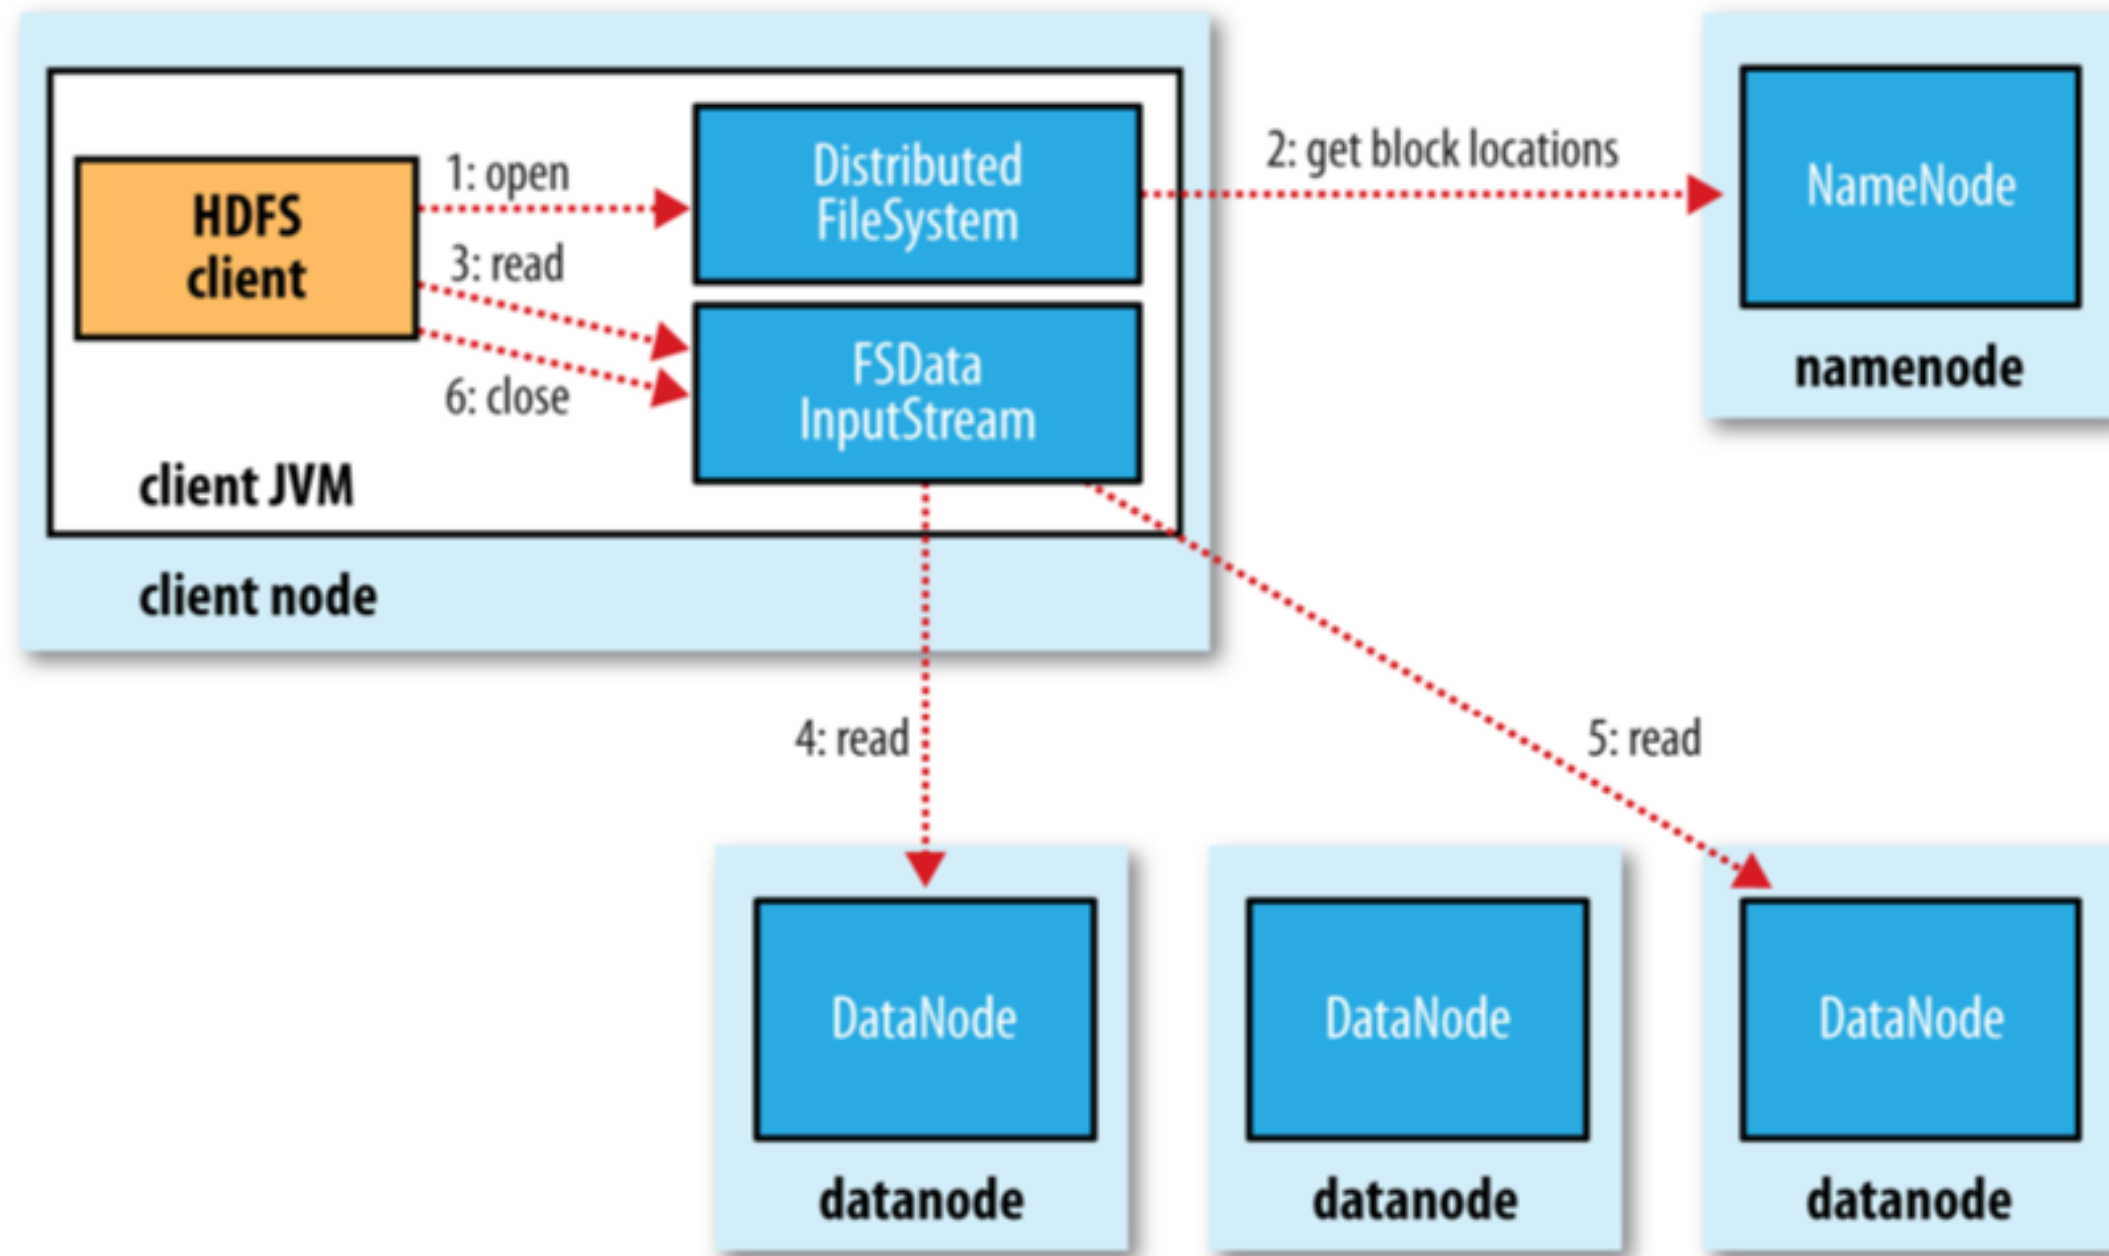


Высокая доступность HDFS

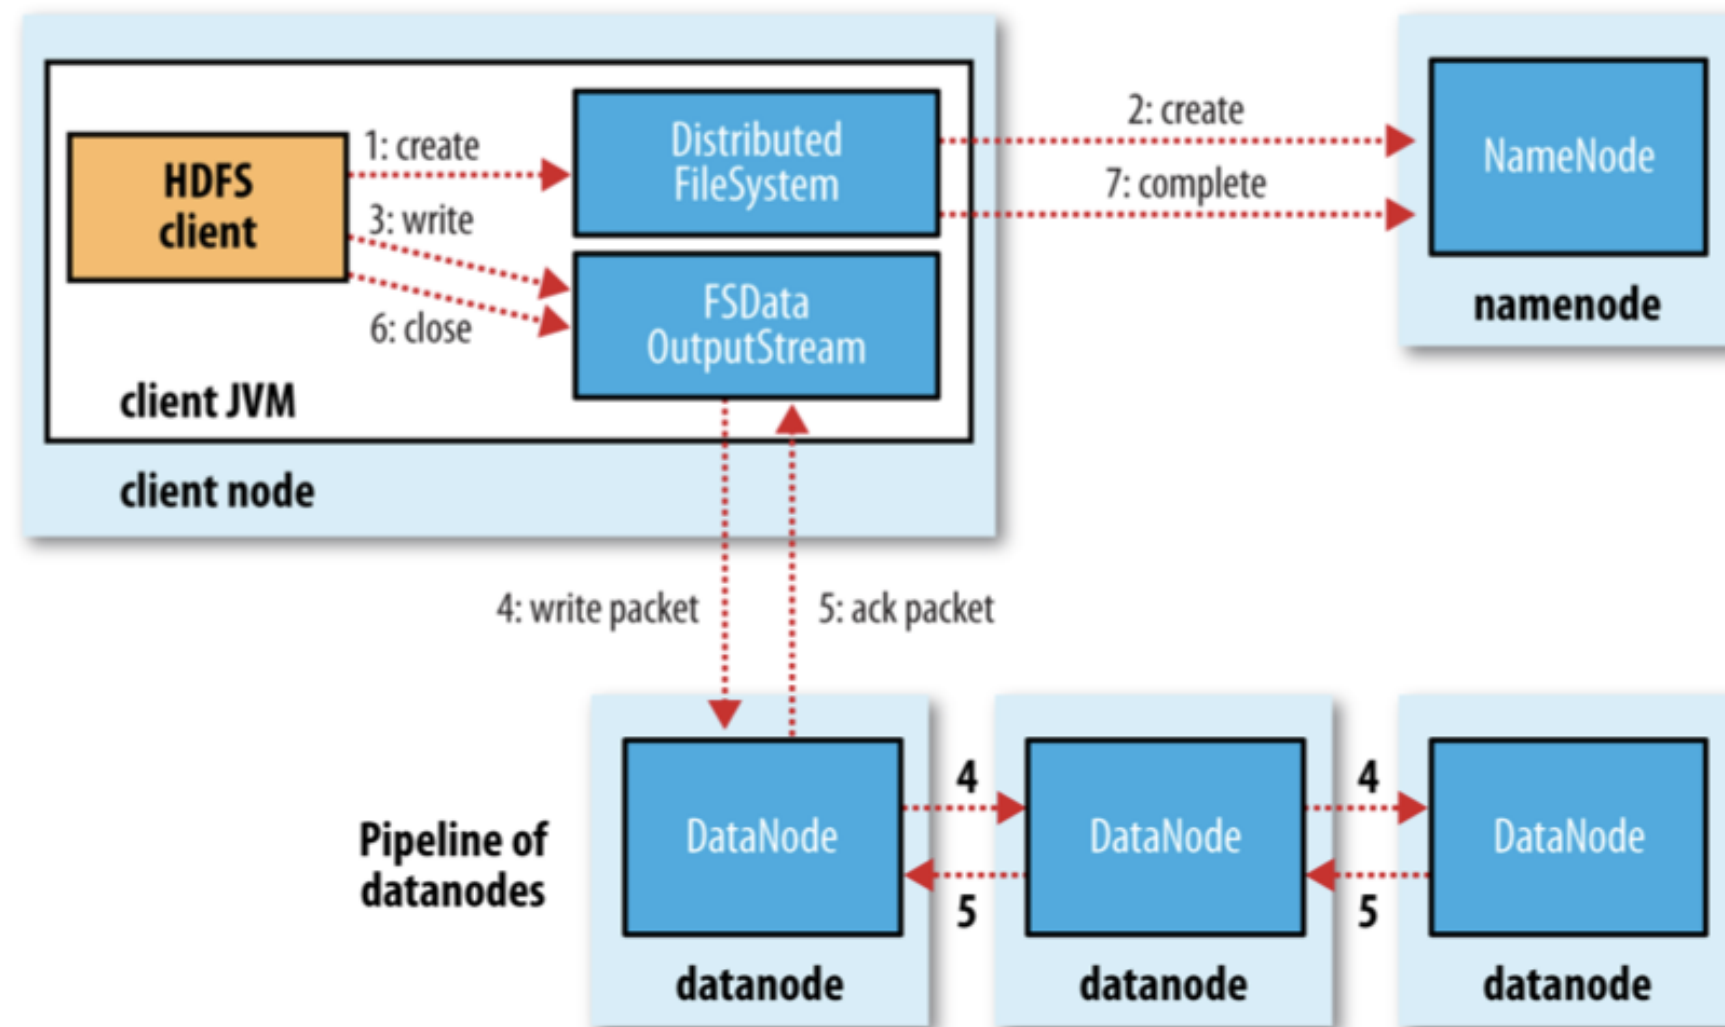
Консольные команды HDFS

- `hadoop distch`
- `hadoop distcp`
- `hadoop fs -put`
- `hadoop fs -get`
- `hadoop fs -appendToFile`
- `hadoop fs -cat`
- `hadoop fs -du`
- И т.д.

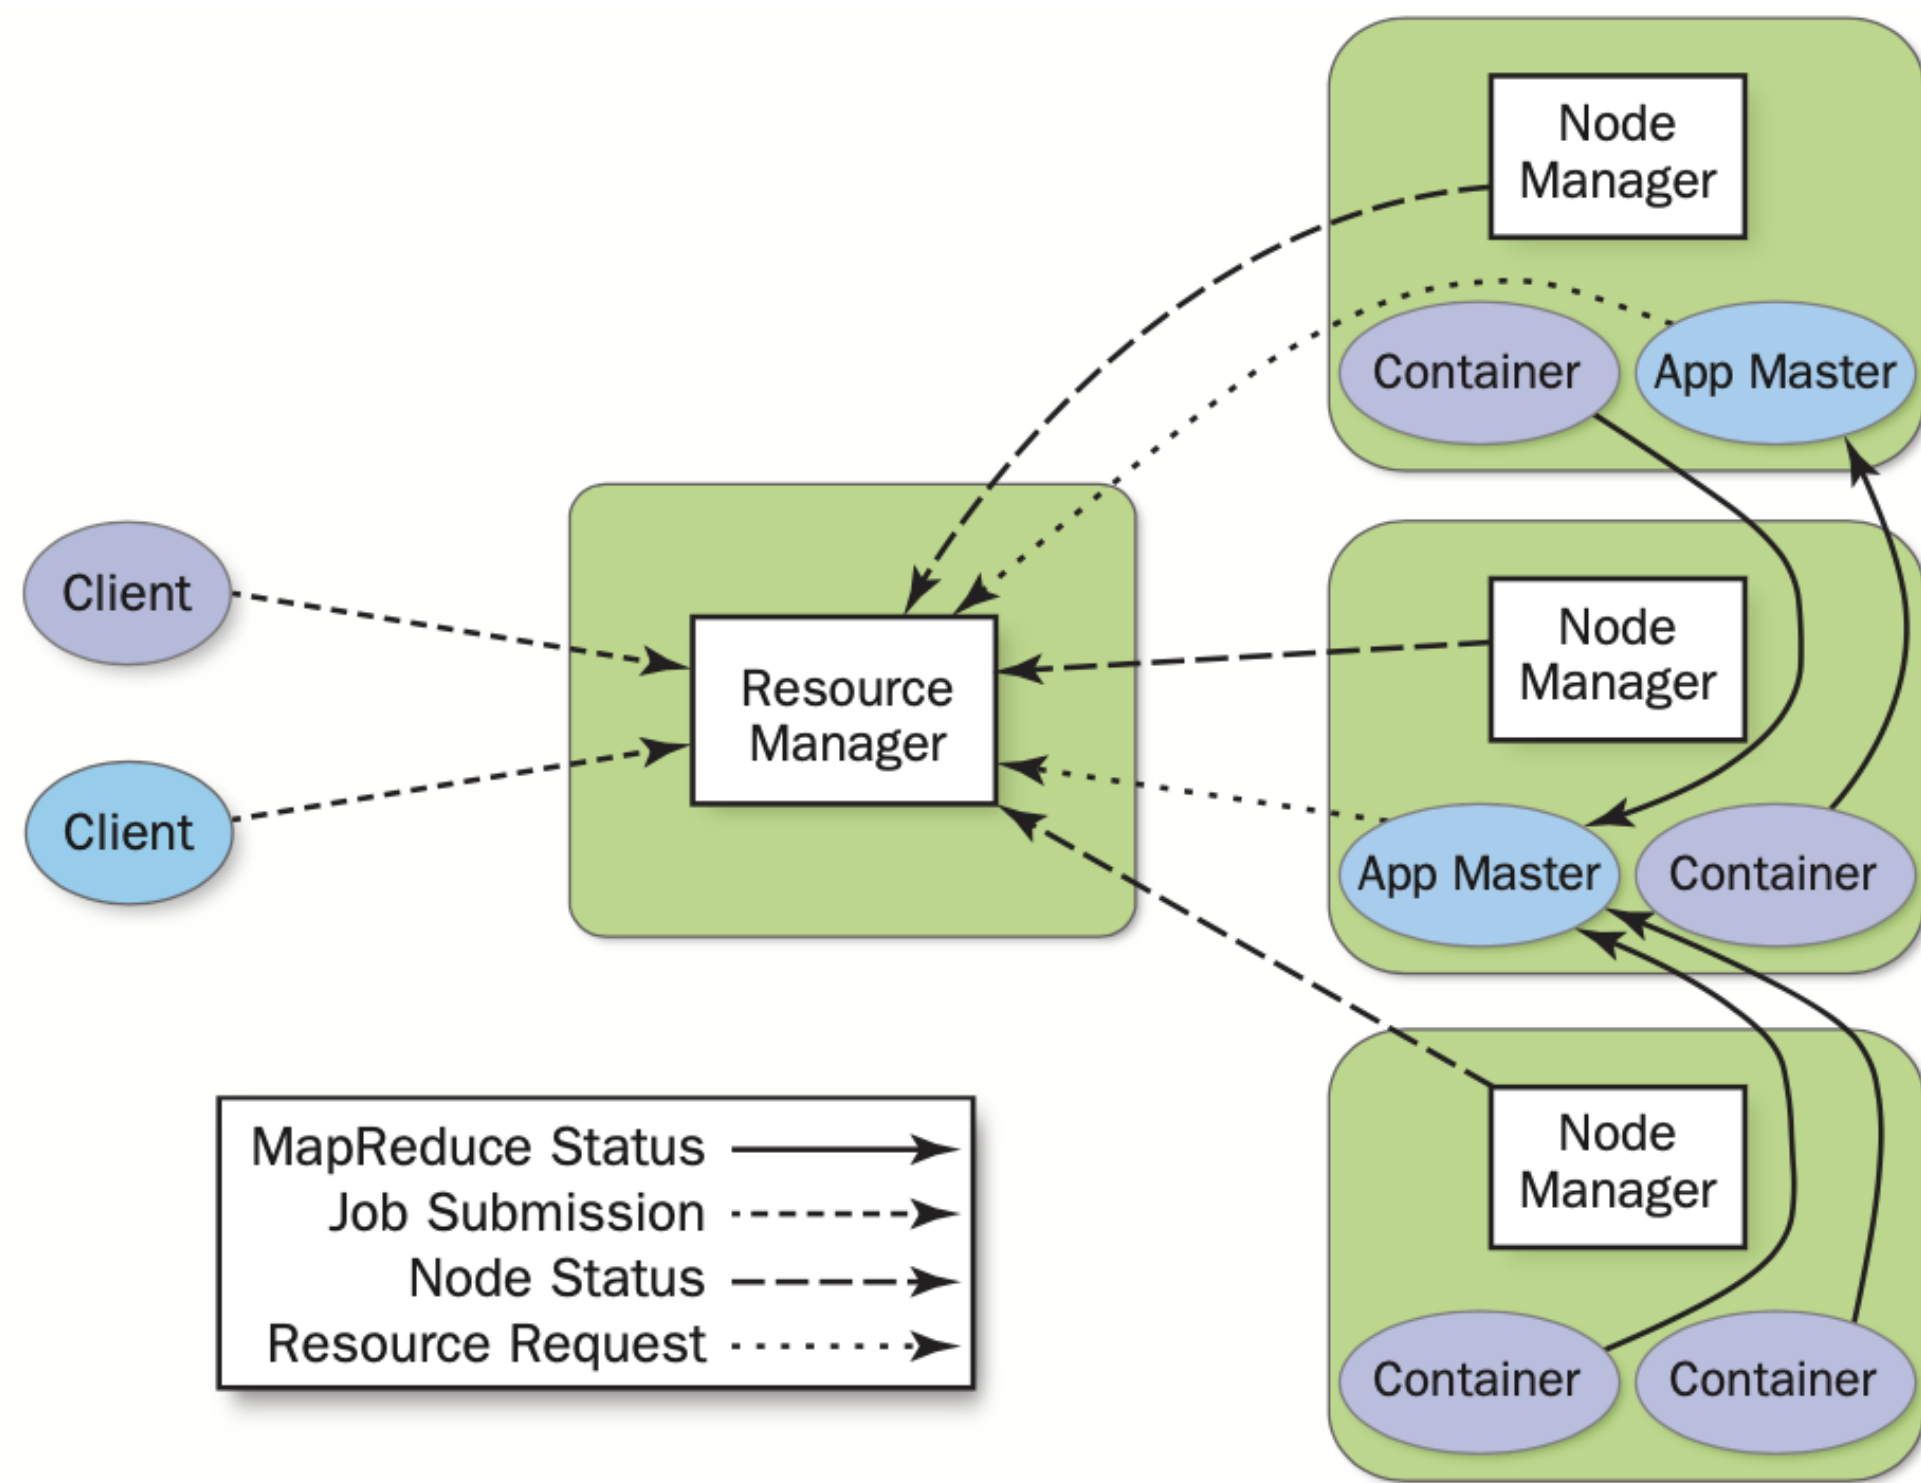
Чтение файла в HDFS



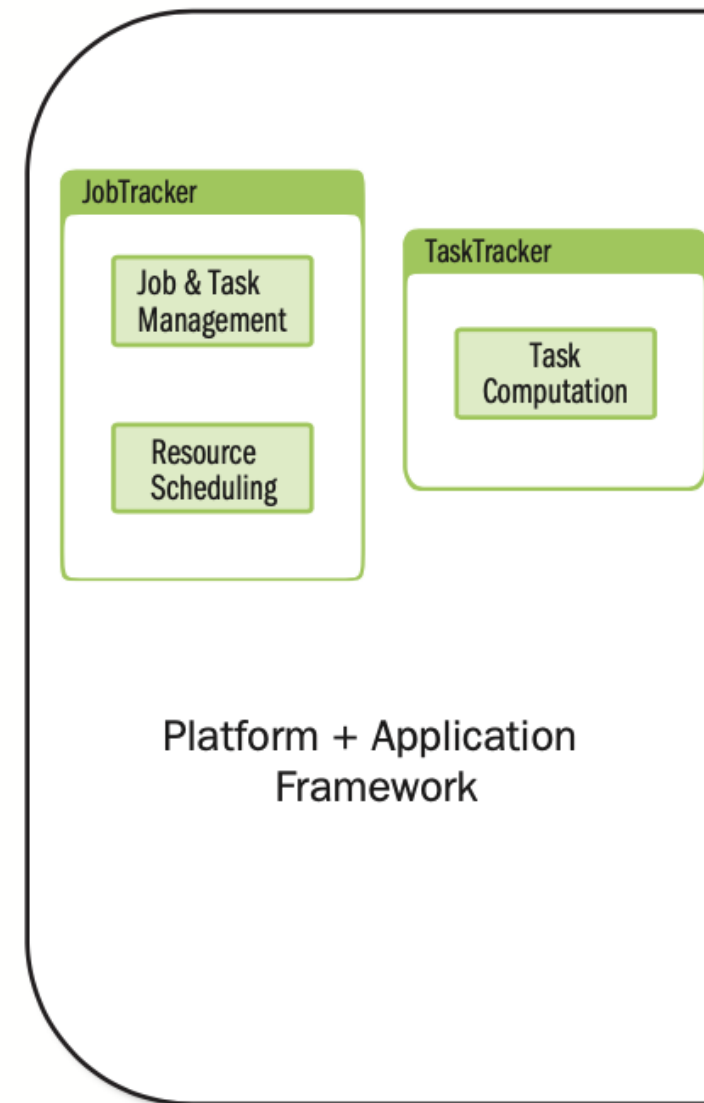
Запись файла в HDFS



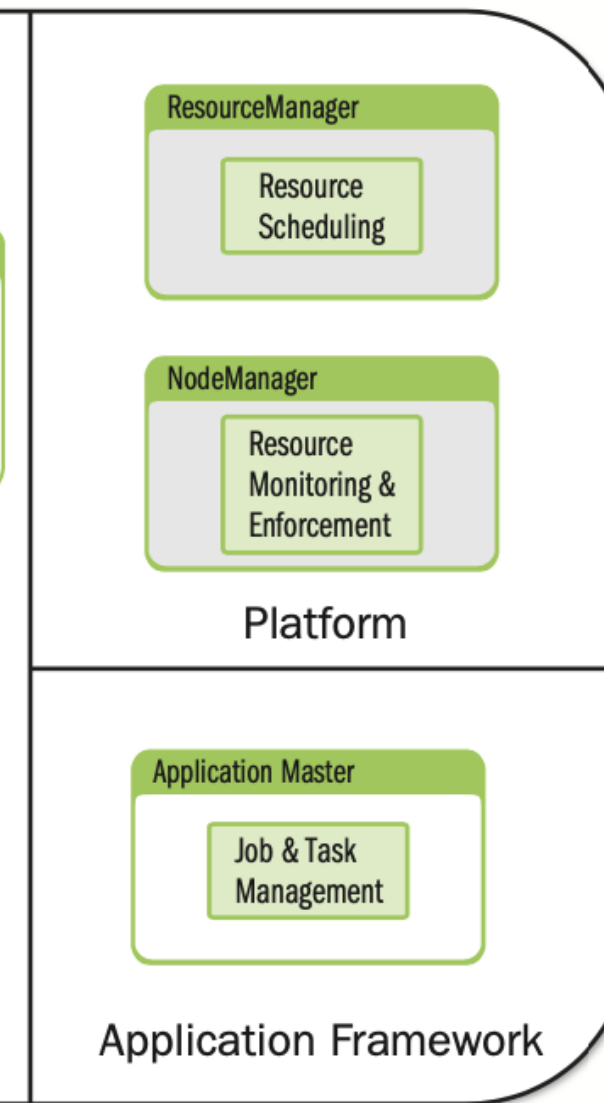
YARN

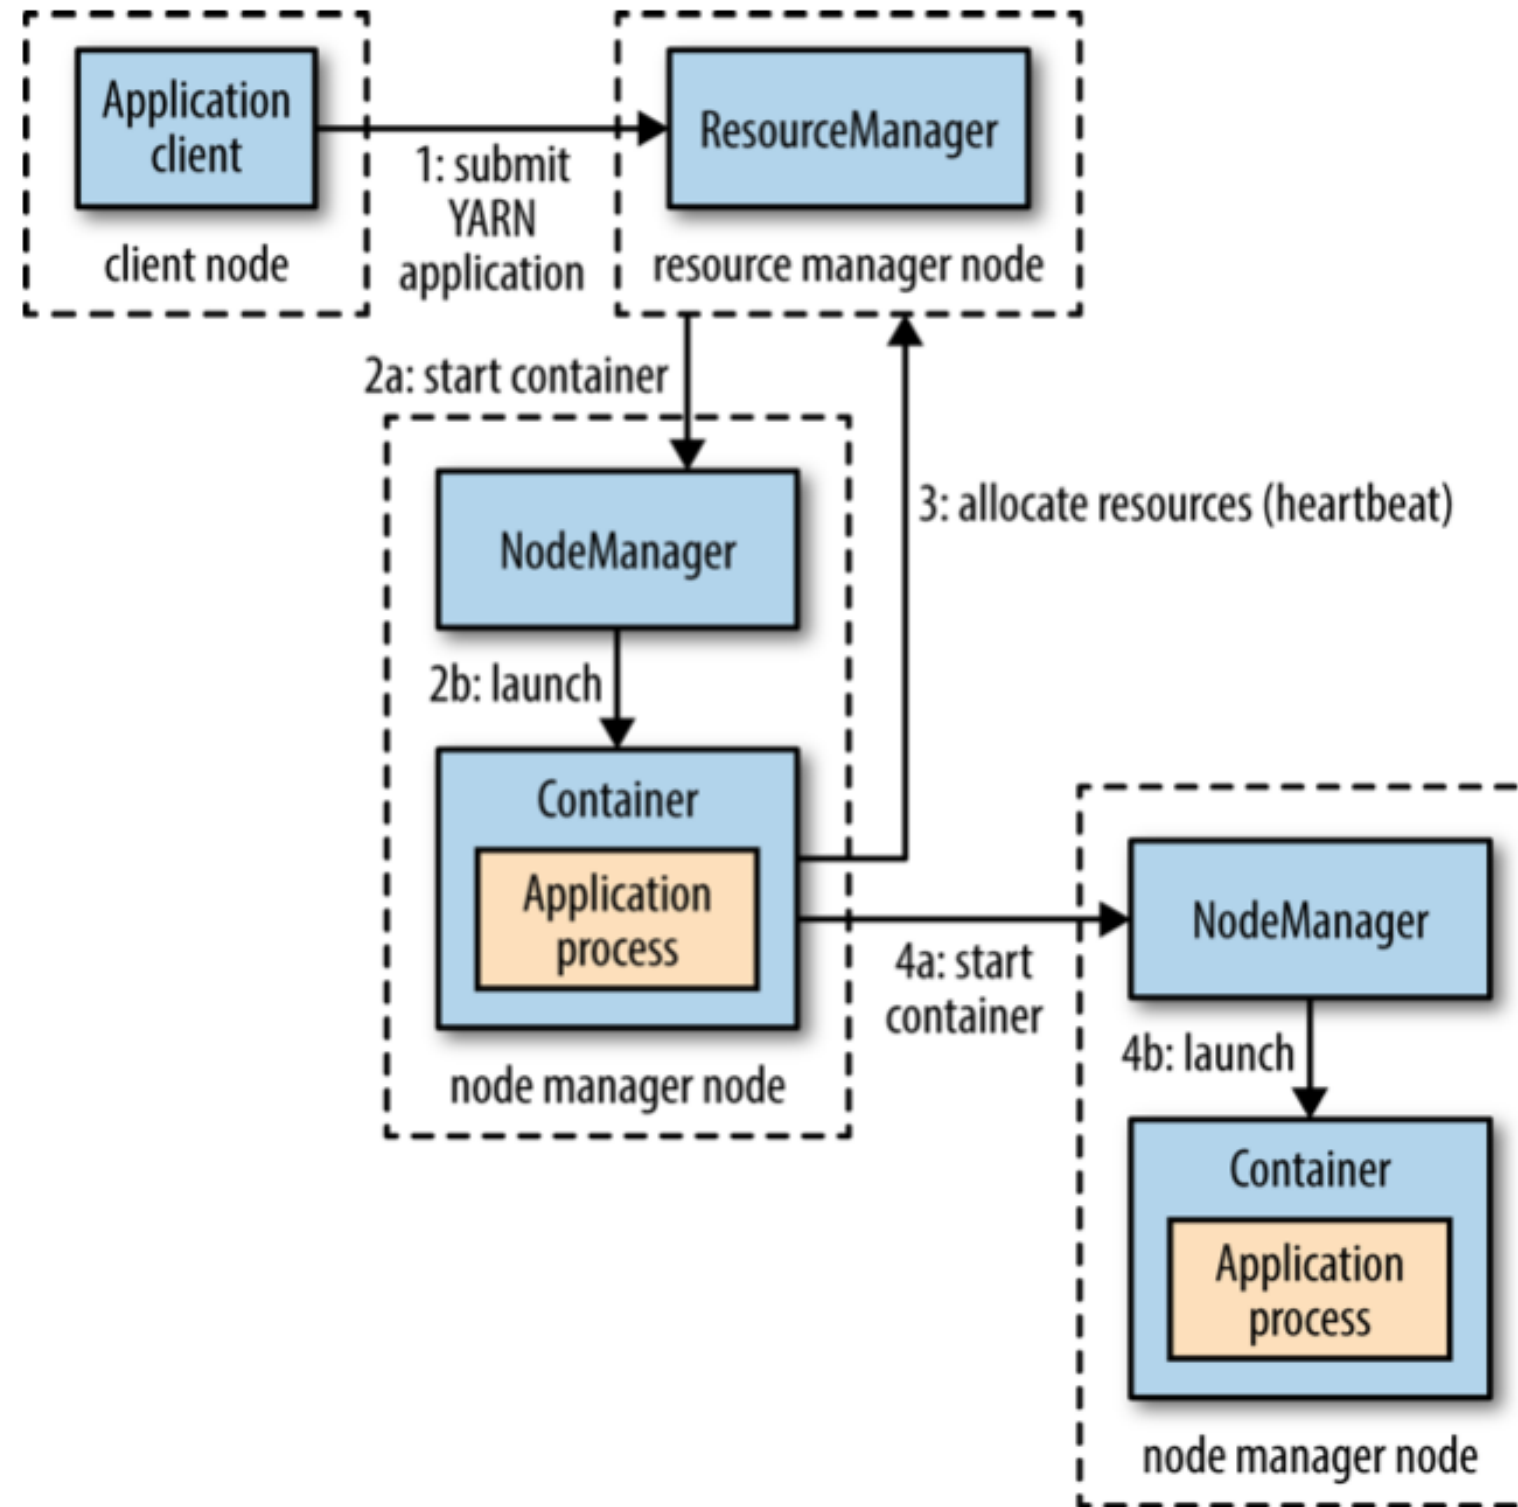


Hadoop 1




Hadoop 2 YARN







i. FIFO Scheduler




ii. Capacity Scheduler

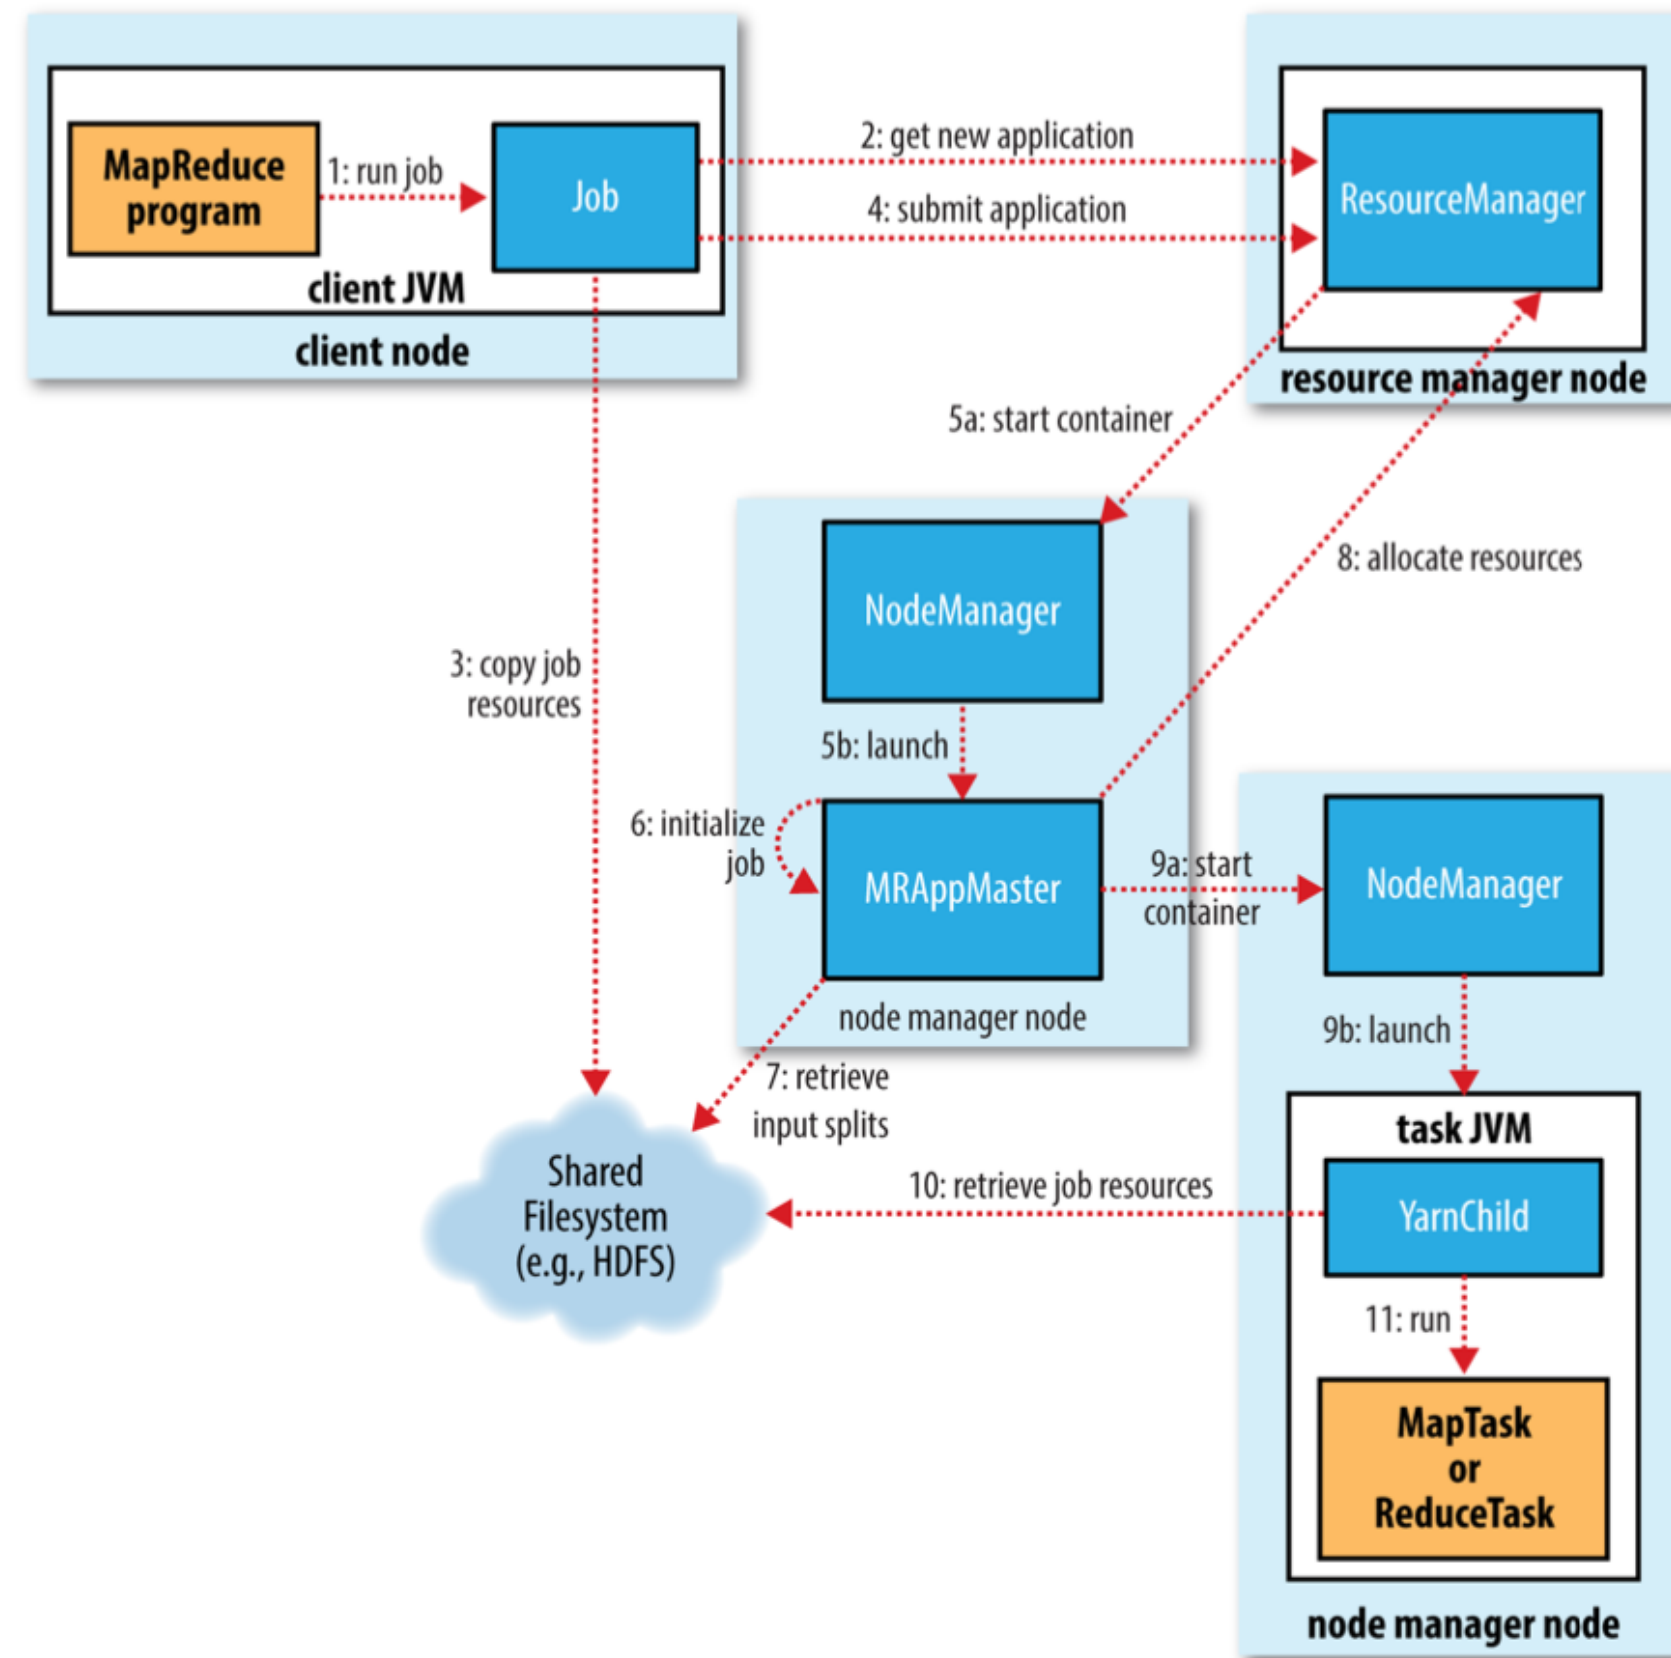


iii. Fair Scheduler





СОБЕРЕМ ВСЕ ВМЕСТЕ



JOIN

Stations

Station ID	Station Name
011990-99999	SIHCCAJAVRI
012650-99999	TYNSET-HANSMOEN

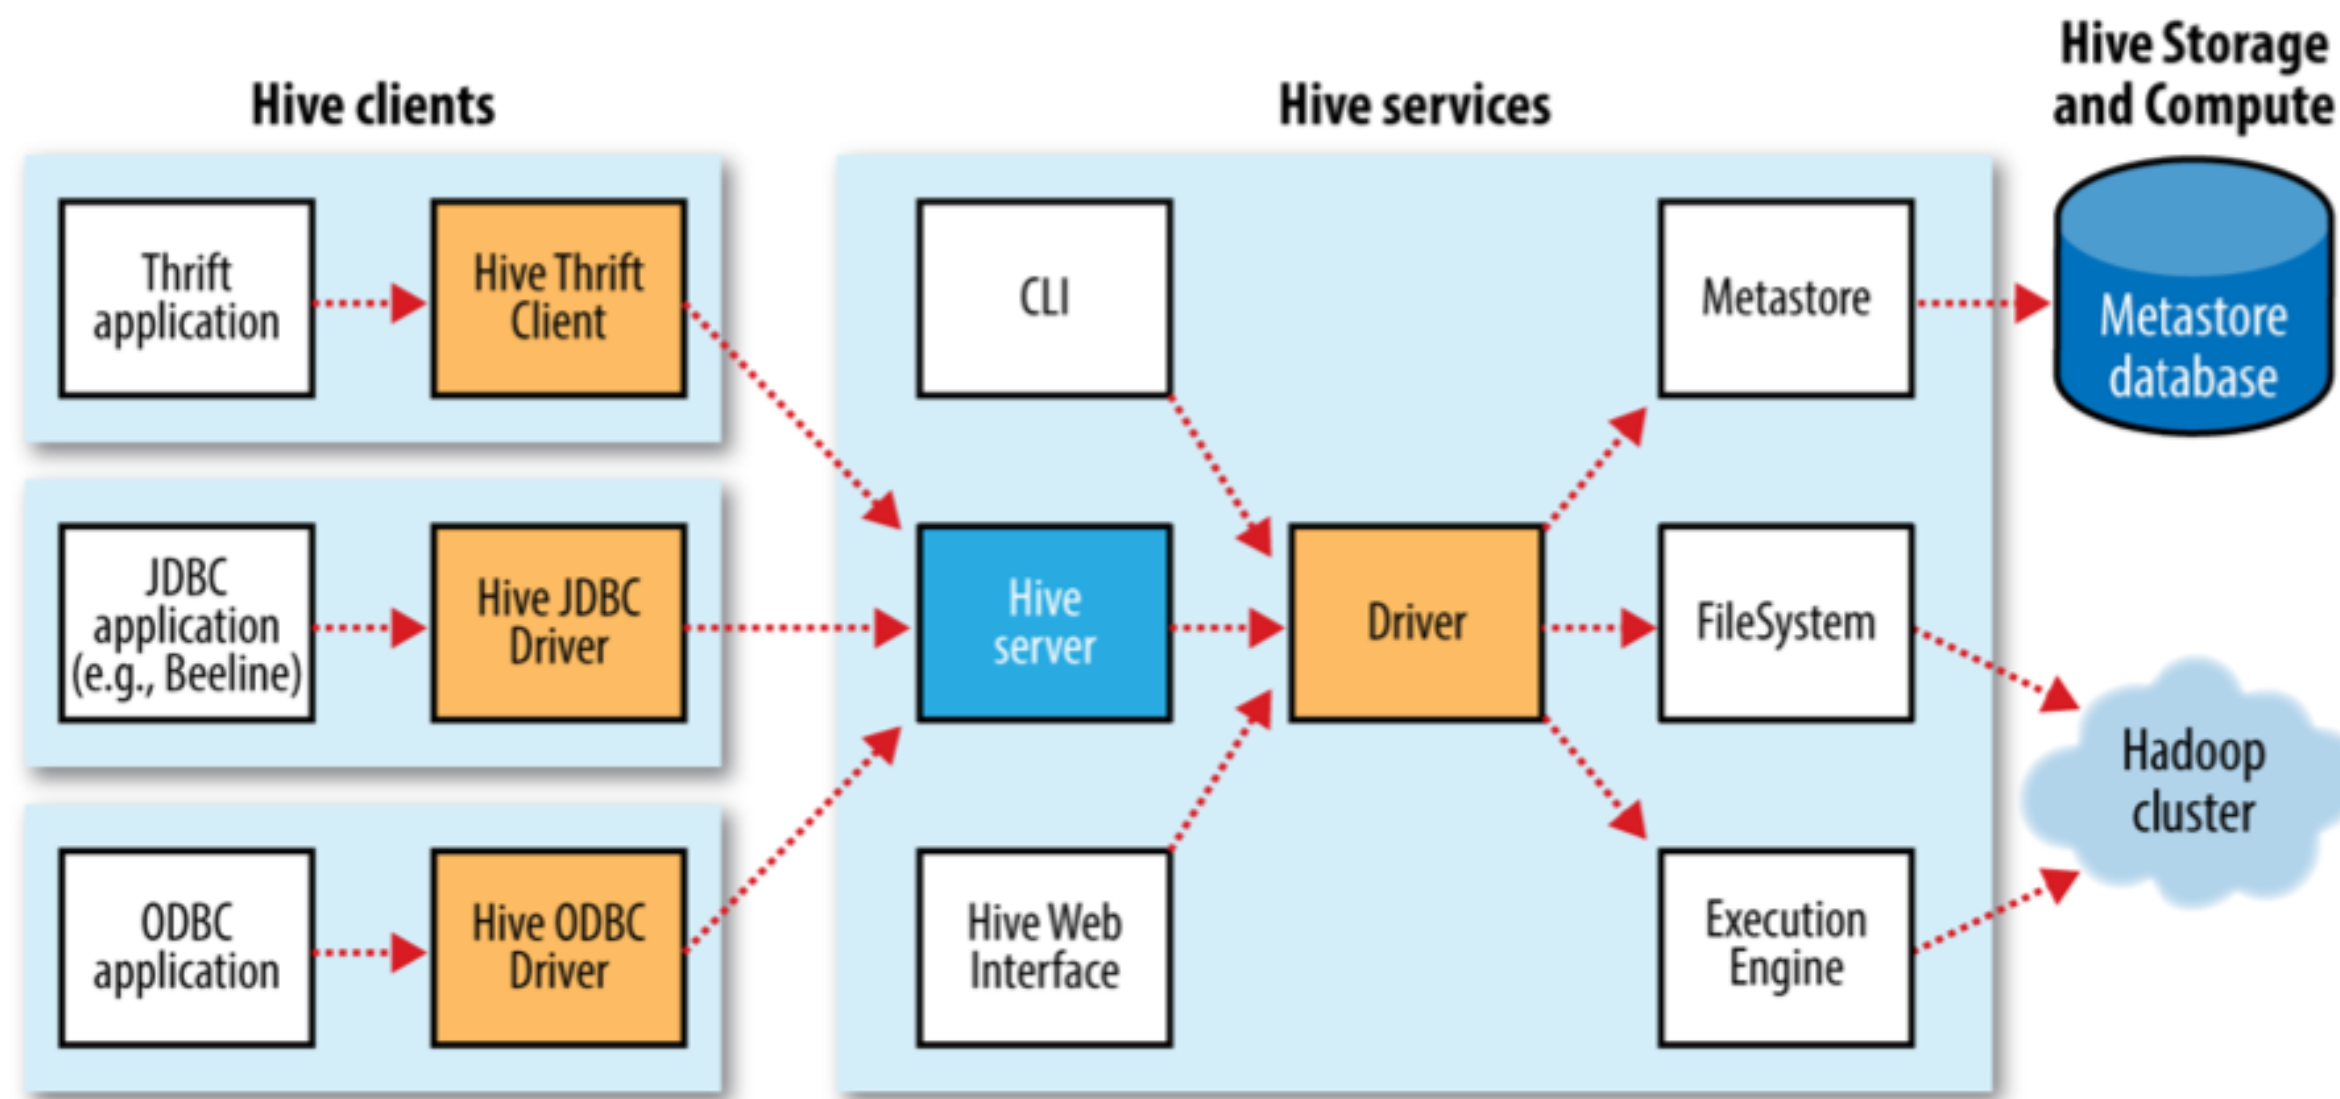
Records

Station ID	Timestamp	Temperature
012650-99999	194903241200	111
012650-99999	194903241800	78
011990-99999	195005150700	0
011990-99999	195005151200	22
011990-99999	195005151800	-11

Join

Station ID	Station Name	Timestamp	Temperature
011990-99999	SIHCCAJAVRI	195005150700	0
011990-99999	SIHCCAJAVRI	195005151200	22
011990-99999	SIHCCAJAVRI	195005151800	-11
012650-99999	TYNSET-HANSMOEN	194903241200	111
012650-99999	TYNSET-HANSMOEN	194903241800	78

HIVE



Managed Tables and External Tables

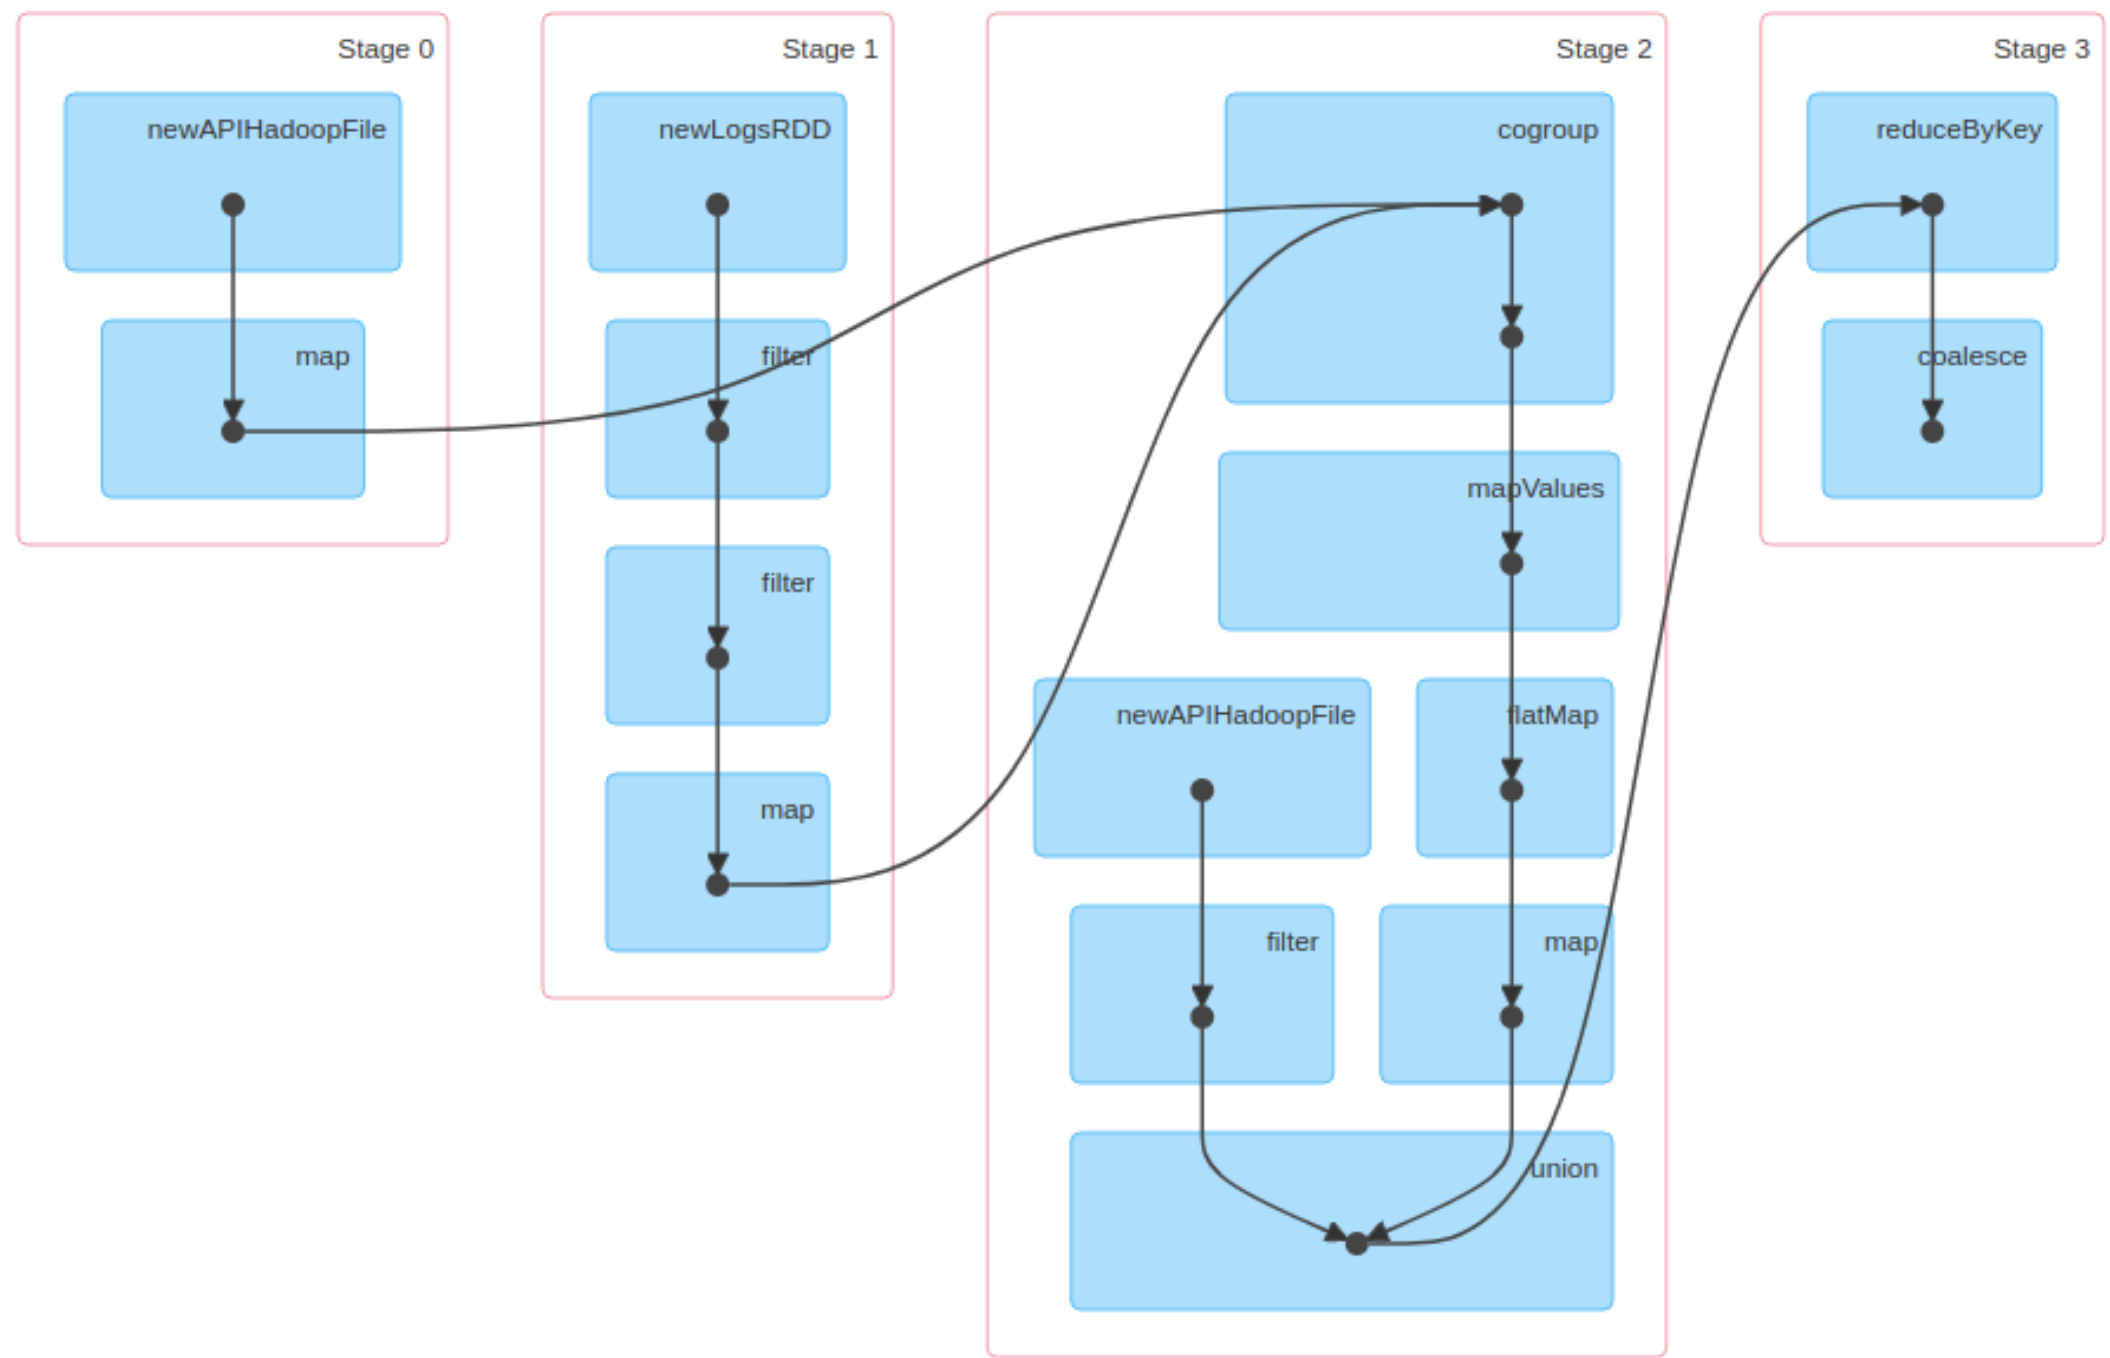
```
CREATE [TEMPORARY] [EXTERNAL] TABLE [IF NOT EXISTS] [db_name.]table_name
[(col_name data_type [column_constraint_specification] [COMMENT col_comment], ... [constraint_specification])]
[COMMENT table_comment]
[PARTITIONED BY (col_name data_type [COMMENT col_comment], ...)]
[CLUSTERED BY (col_name, col_name, ...) [SORTED BY (col_name [ASC | DESC], ...)] INTO num_buckets BUCKETS]
[ROW FORMAT row_format]
[STORED AS file_format]
  | STORED BY 'storage.handler.class.name' [WITH SERDEPROPERTIES (...)] -- (Note: Available in Hive 0.6.0 and later)
]
[LOCATION hdfs_path]
[AS select_statement];
```

*указано не все!

JOIN

```
SELECT /*+ MAPJOIN(b) */ a.key, a.value  
FROM a JOIN b ON a.key = b.key
```

SPARK



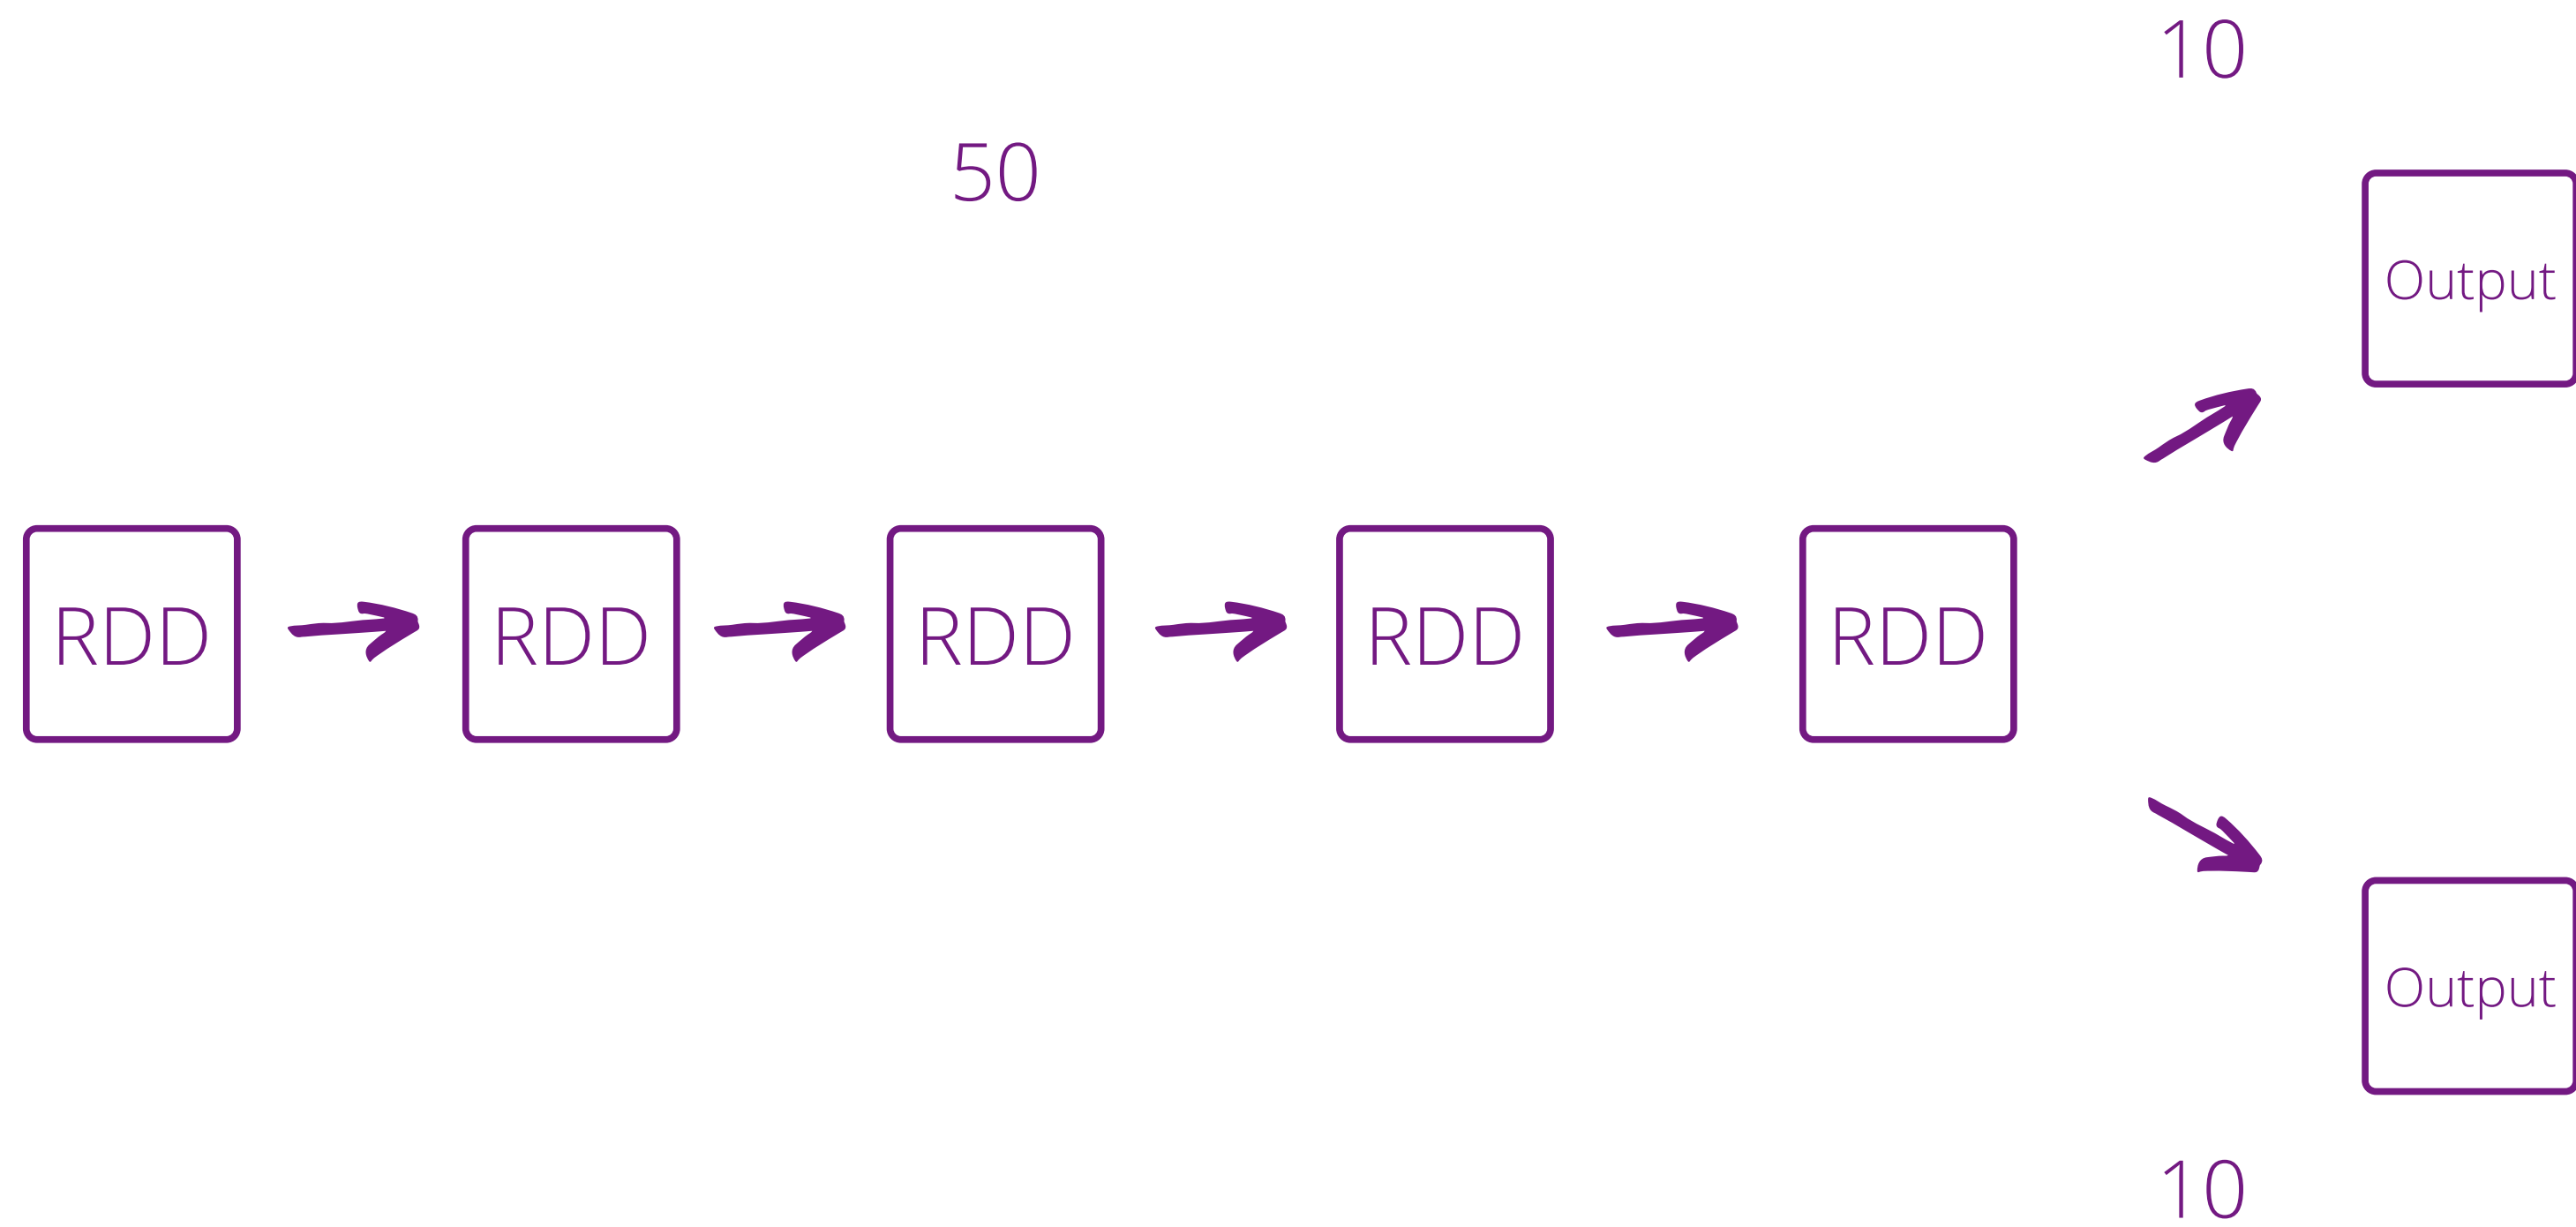
Resilient Distributed Datasets (RDD)

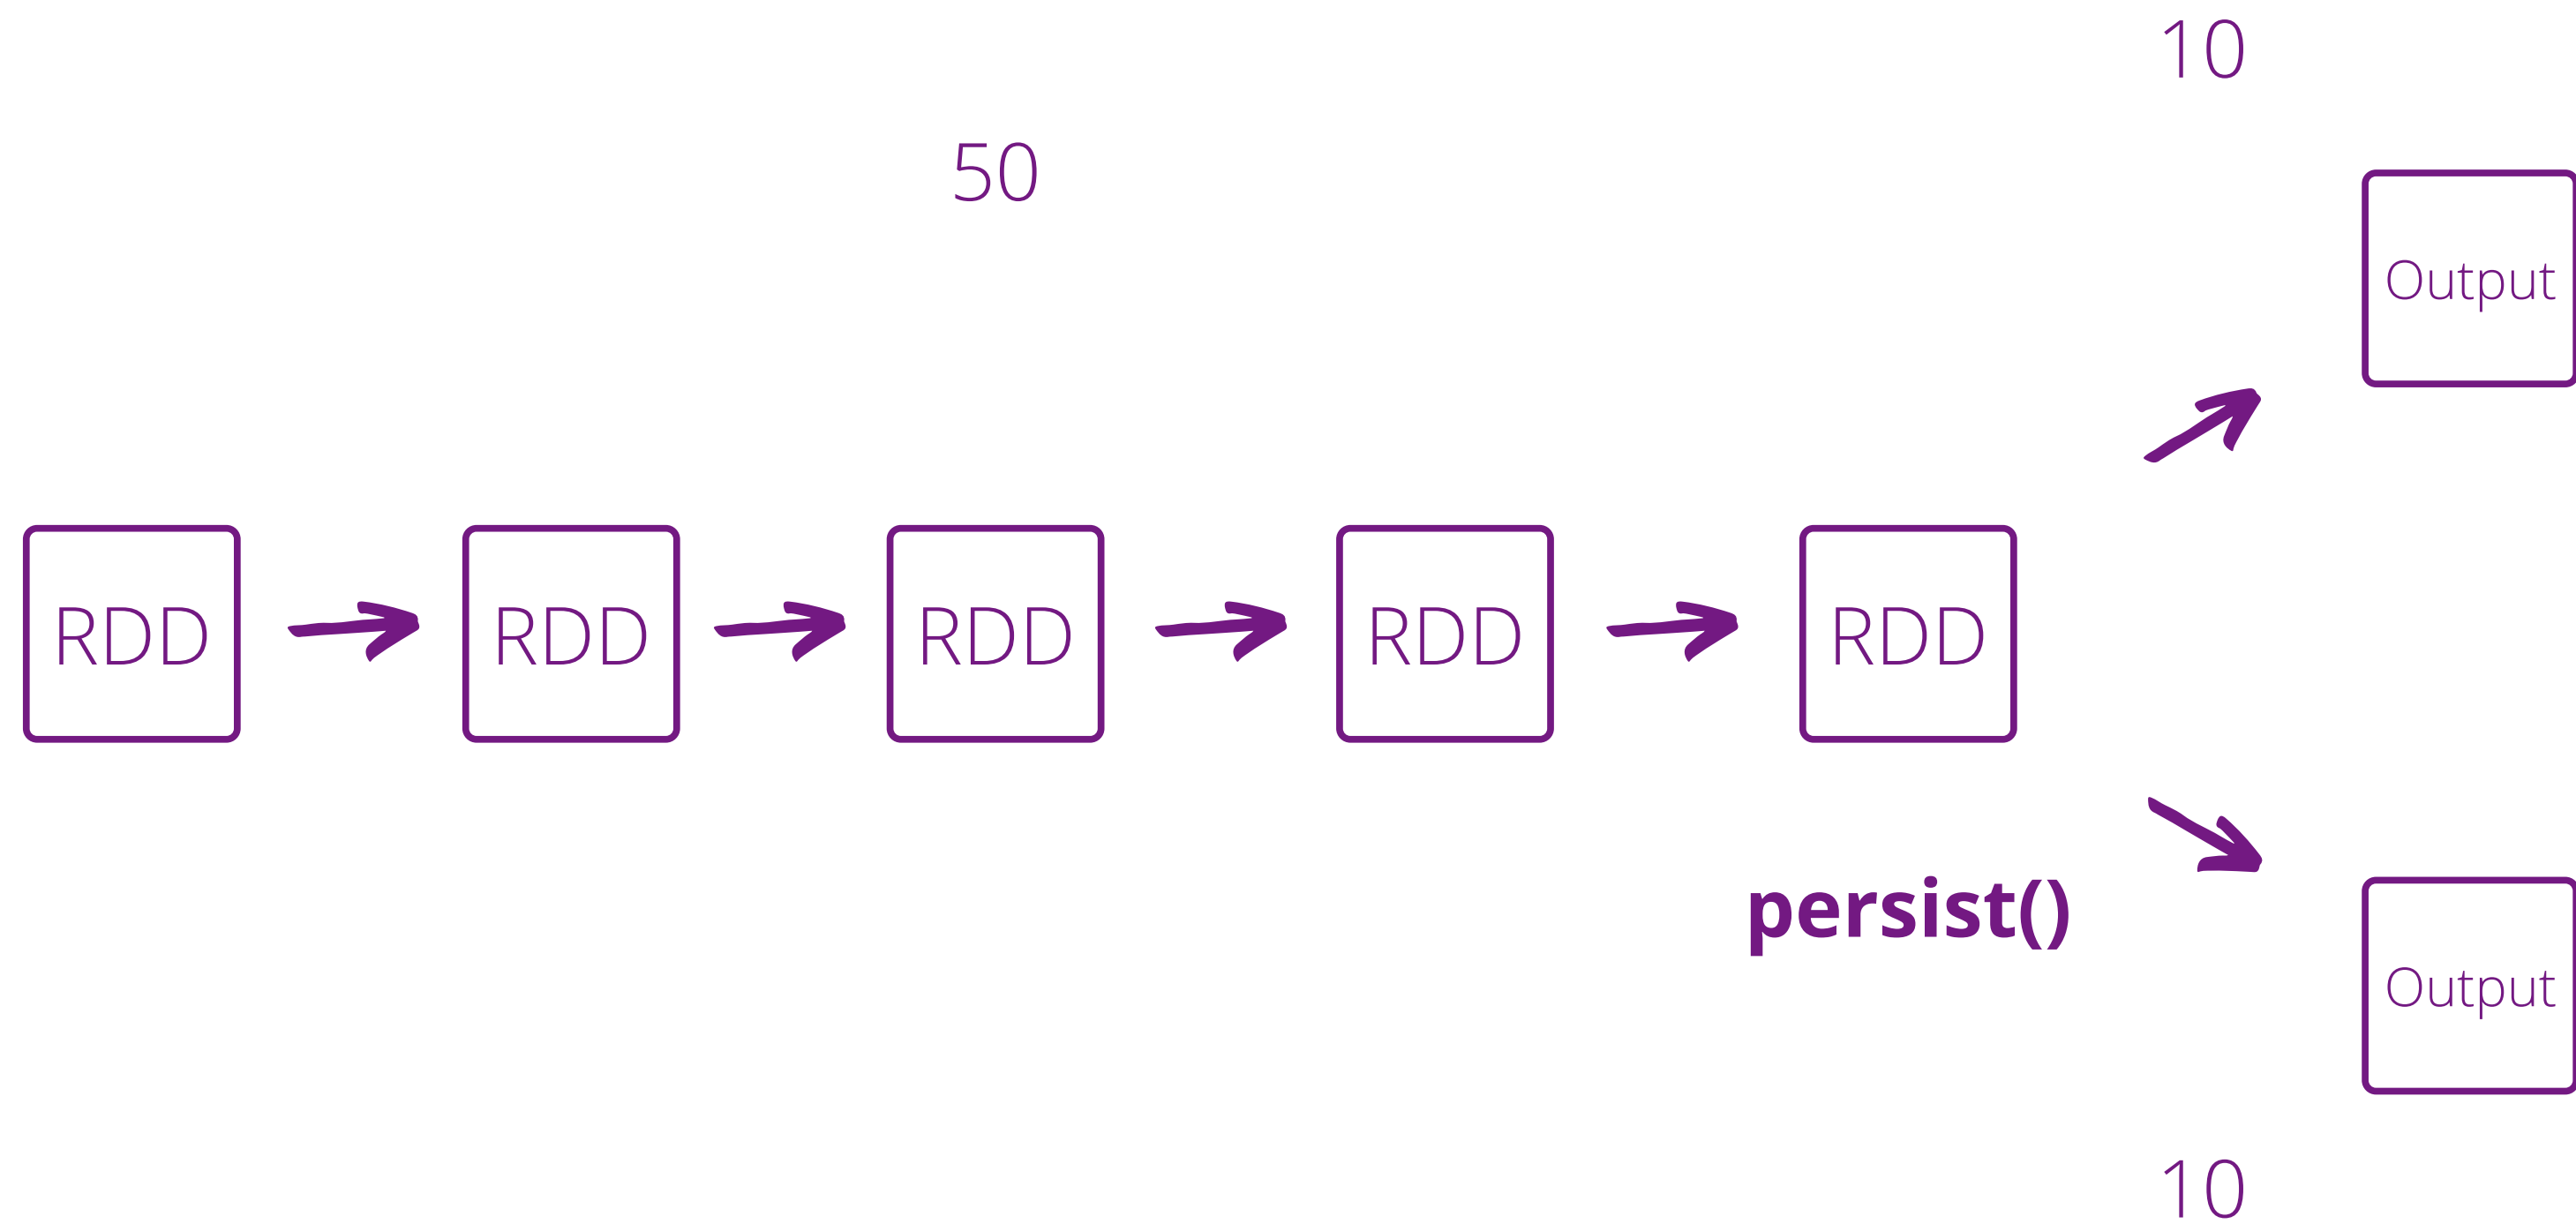
Transformations and Actions

- `map(func)`
- `filter(func)`
- `flatMap(func)`
- `sample(withReplacement, fraction, seed)`
- `union(otherDataset)`
- `distinct([numPartitions])`
- `reduceByKey(func, [numPartitions])`
- `join(otherDataset, [numPartitions])`
- `repartition(numPartitions)`
- `coalesce(numPartitions)`

- `collect()`
- `count()`
- `take(n)`
- `takeSample(withReplacement, num, [seed])`
- `saveAsTextFile(path)`

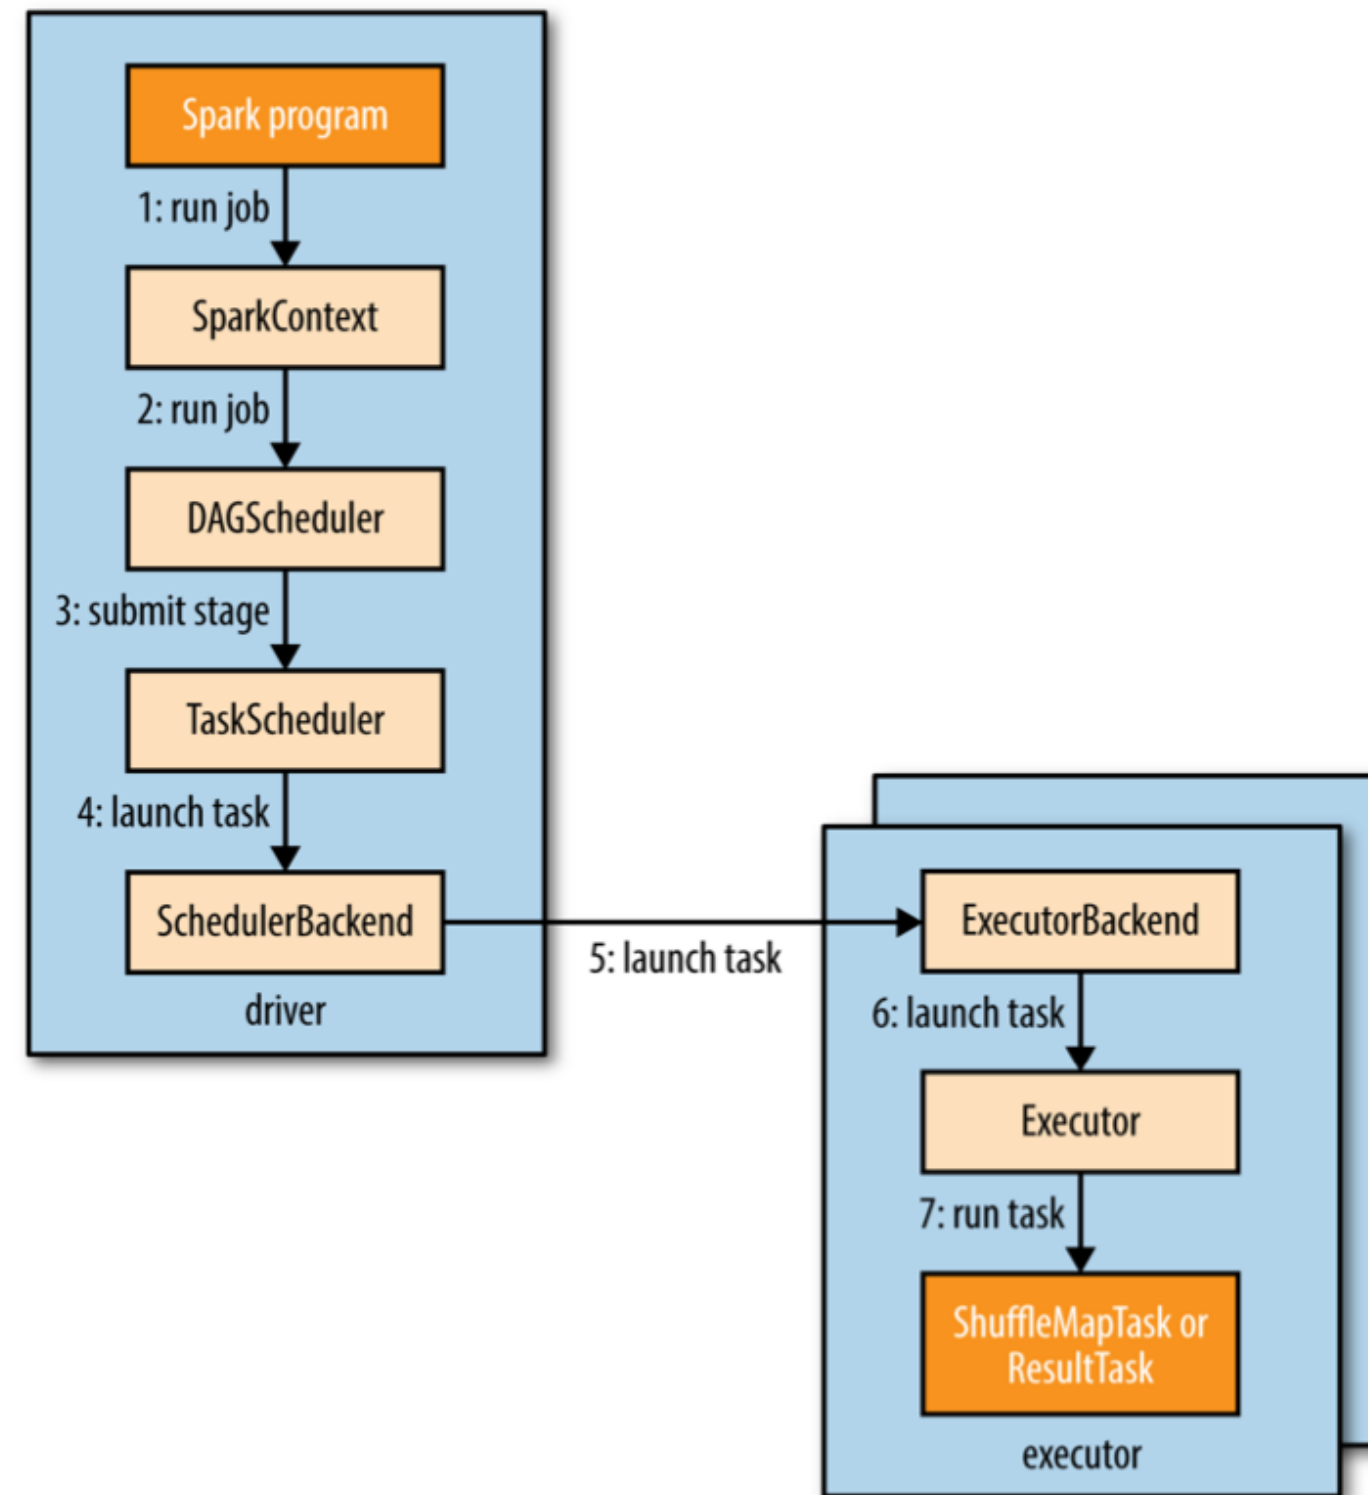
```
lines = sc.textFile("data.txt")  
pairs = lines.map(lambda s: (s, 1))  
counts = pairs.reduceByKey(lambda a, b: a + b)
```

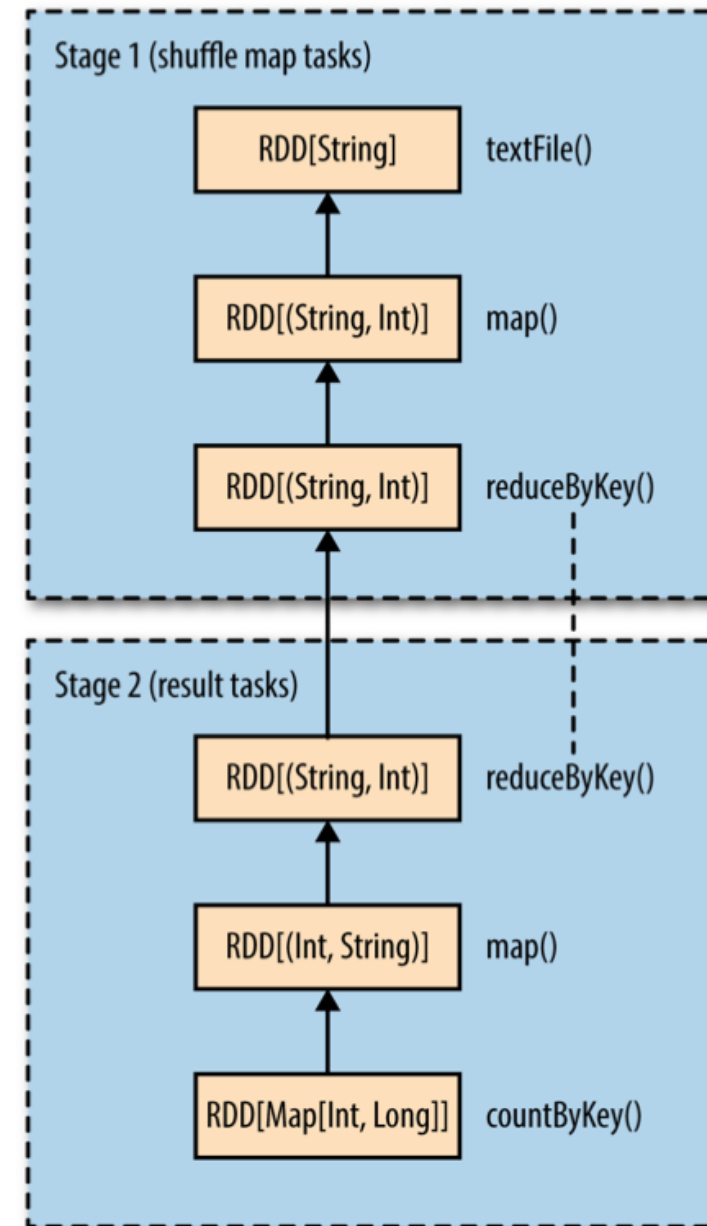




- MEMORY_ONLY
- MEMORY_AND_DISK
- MEMORY_ONLY_SER (Java and Scala)
- MEMORY_AND_DISK_SER (Java and Scala)
- DISK_ONLY
- MEMORY_ONLY_2
- MEMORY_AND_DISK_2

```
sc.broadcast([1, 2, 3])
```



→ RDD dependency

Shuffle

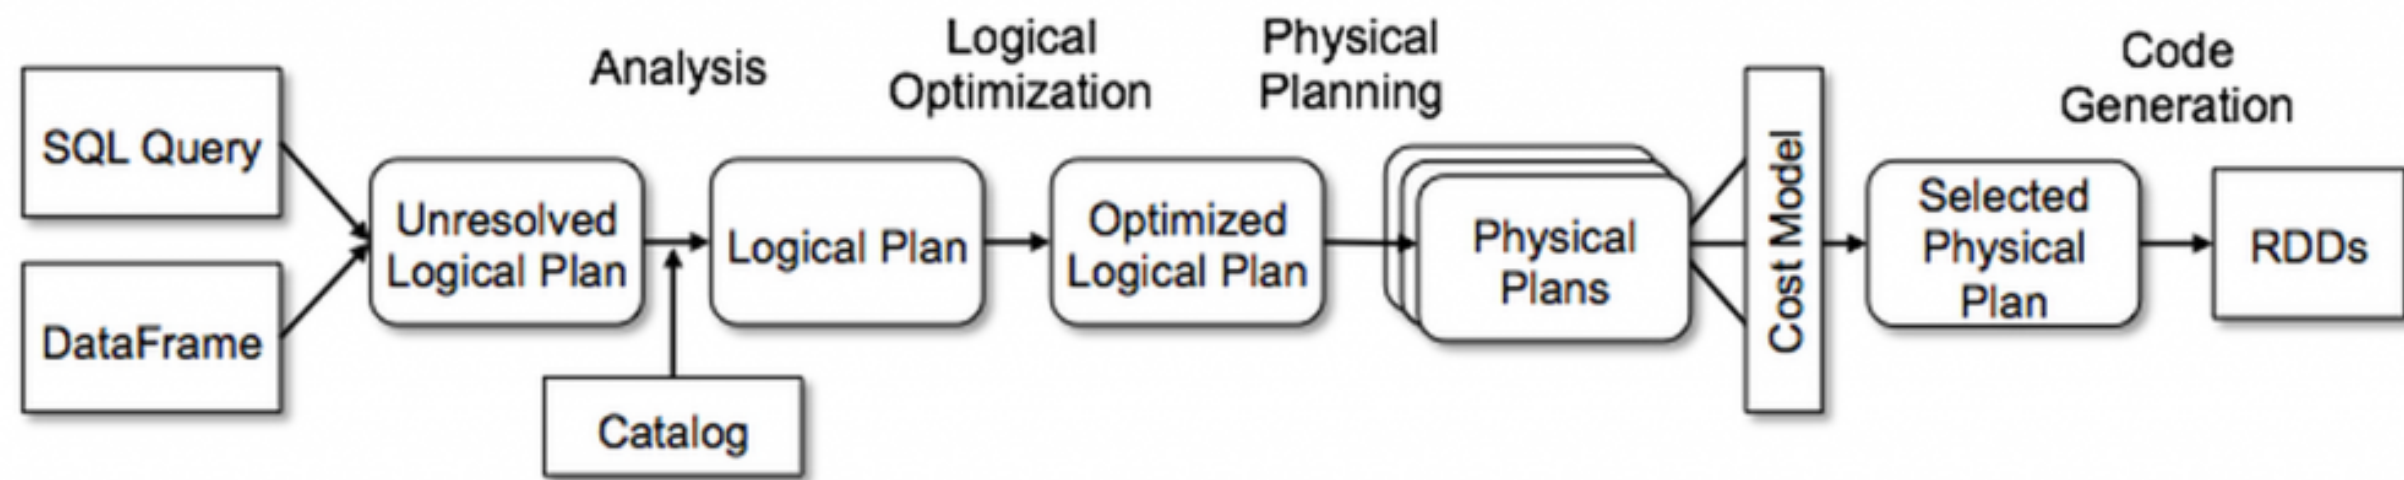
DataFrame



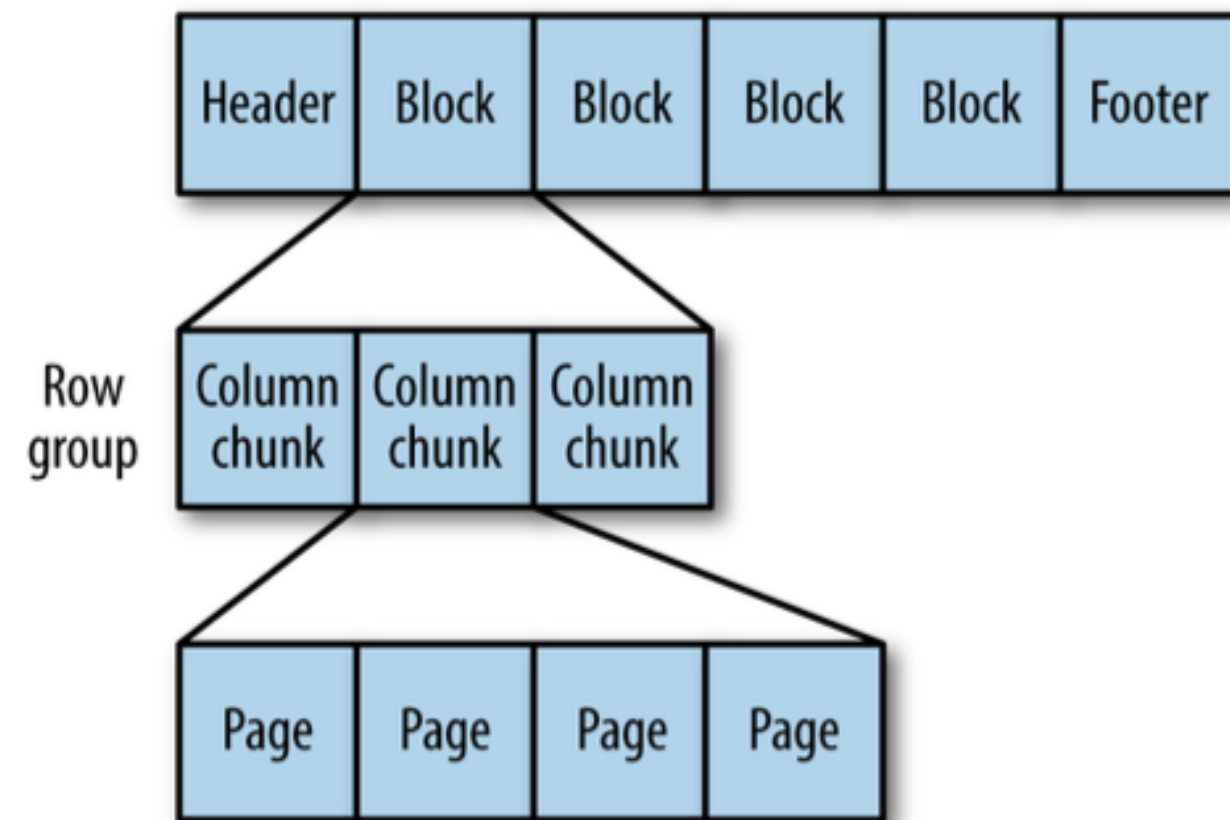
- `df = spark.read.json("examples/src/main/resources/people.json")`
- `df.filter(df['age'] > 21).show()`
- `df.groupBy("age").count().show()`
- `sqlDF = spark.sql("SELECT * FROM people")`

```
jdbcDF = spark.read \  
  .format("jdbc") \  
  .option("url", "jdbc:postgresql:dbserver") \  
  .option("dbtable", "schema.tablename") \  
  .option("user", "username") \  
  .option("password", "password") \  
  .load()
```

- `df = spark.read.parquet("examples/src/main/resources/users.parquet")`
- `(df .write .partitionBy("favorite_color") .bucketBy(42, "name") .saveAsTable("people_partitioned_bucketed"))`



PARQUET





@almorozovv

Thank you!